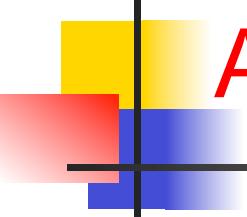


Content-based Image Retrieval (CBIR)

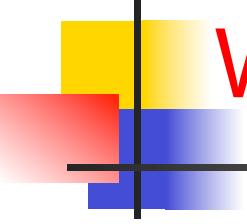
Searching a large database for images
that *match* a query:

- What kinds of databases?
- What kinds of queries?
- What constitutes a match?
- How do we make such searches efficient?



Applications

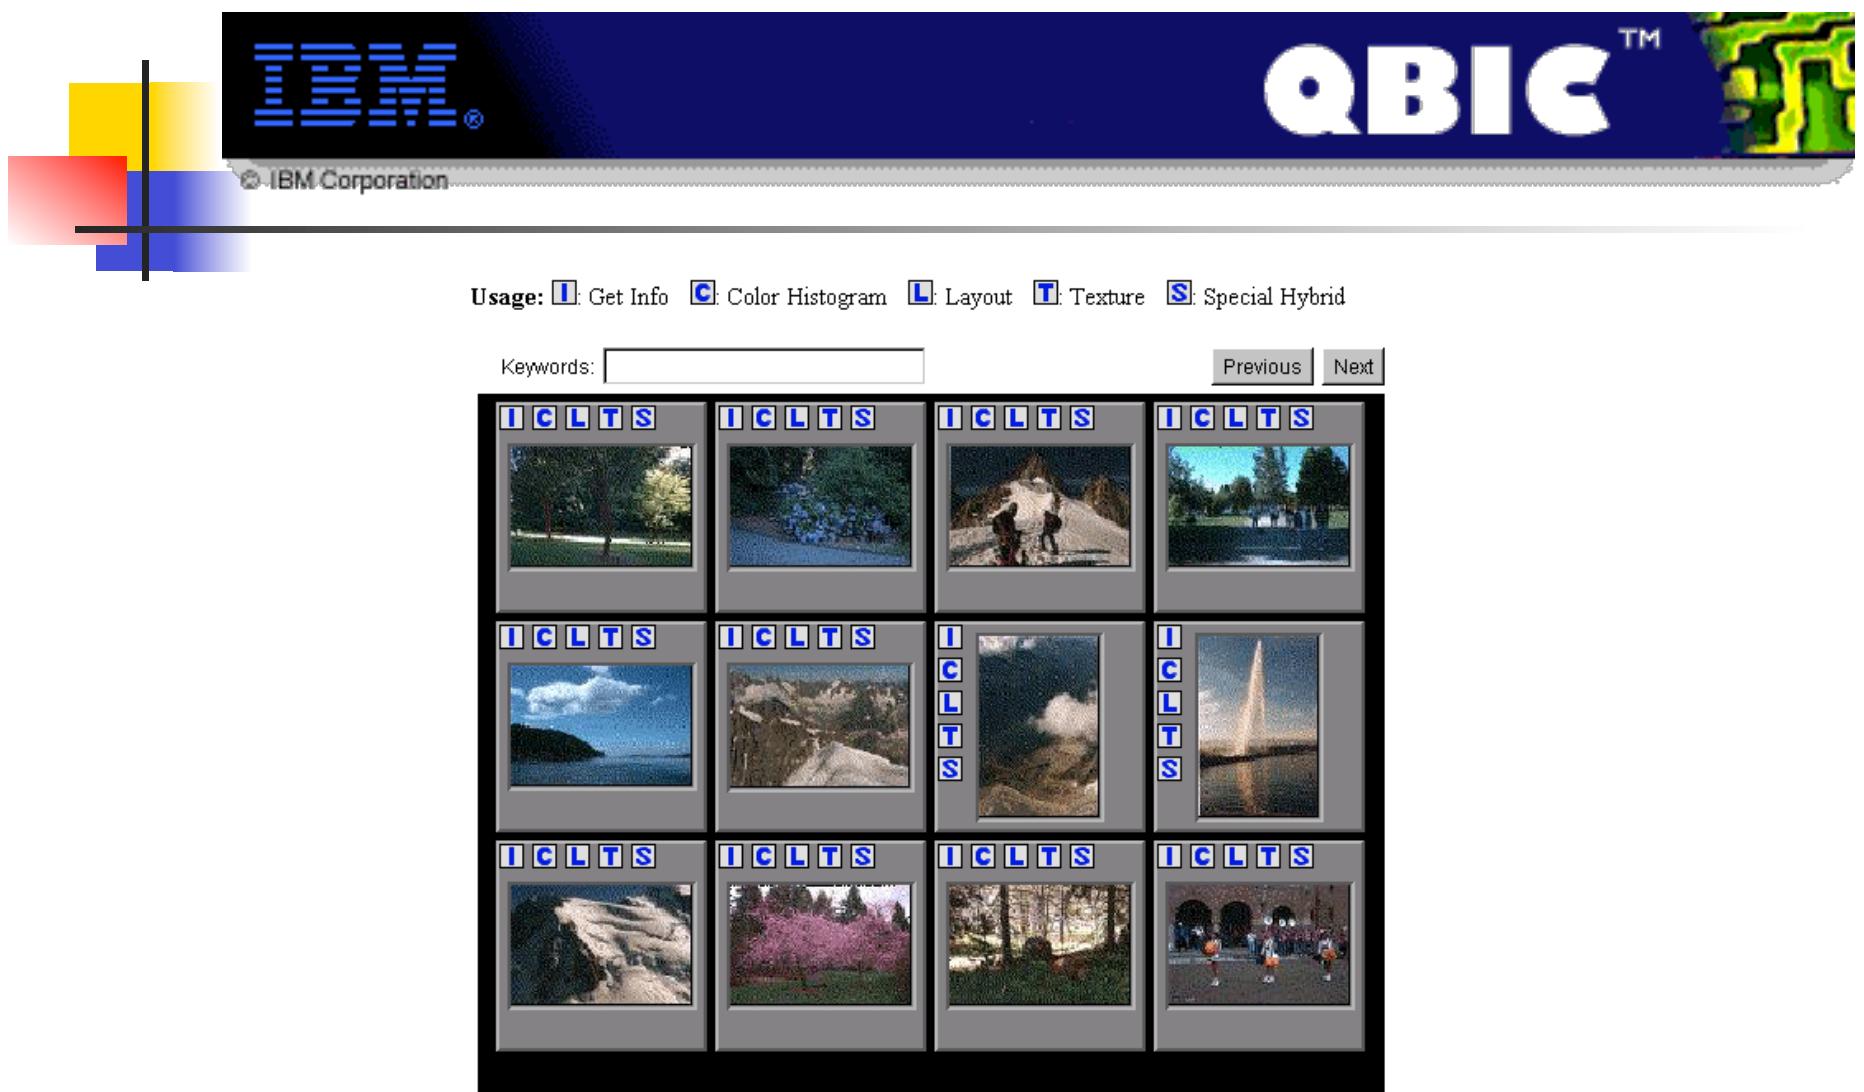
- Art Collections
 - e.g. Fine Arts Museum of San Francisco
- Medical Image Databases
 - CT, MRI, Ultrasound, The Visible Human
- Scientific Databases
 - e.g. Earth Sciences
- General Image Collections for Licensing
 - Corbis, Getty Images
- The World Wide Web



What is a query?

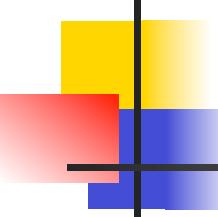
- an **image** you already have
- a rough **sketch** you draw
- a **symbolic description** of what you want
 - e.g. an image of a man and a woman on a beach,
an image of a dad and a daughter on a beach

SYSTEMS



Query was:

Random



Some Systems You Can Try

Corbis Stock Photography and Pictures

<http://www.corbisimages.com>

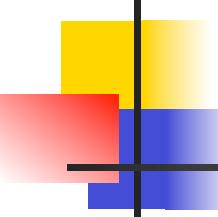
- Corbis sells high-quality images for use in advertising, marketing, illustrating, etc.
- Search is entirely by keywords.
- Human indexers look at each new image and enter keywords.
- A thesaurus constructed from user queries is used.



IBM's QBIC (Query by Image Content)

<http://www.research.ibm.com/topics/popups/deep/manage/html/qbic.html>

- The first commercial system.
- Uses or has-used color percentages, color layout, texture, shape, location, and keywords.



Informedia

CMU's informedia project

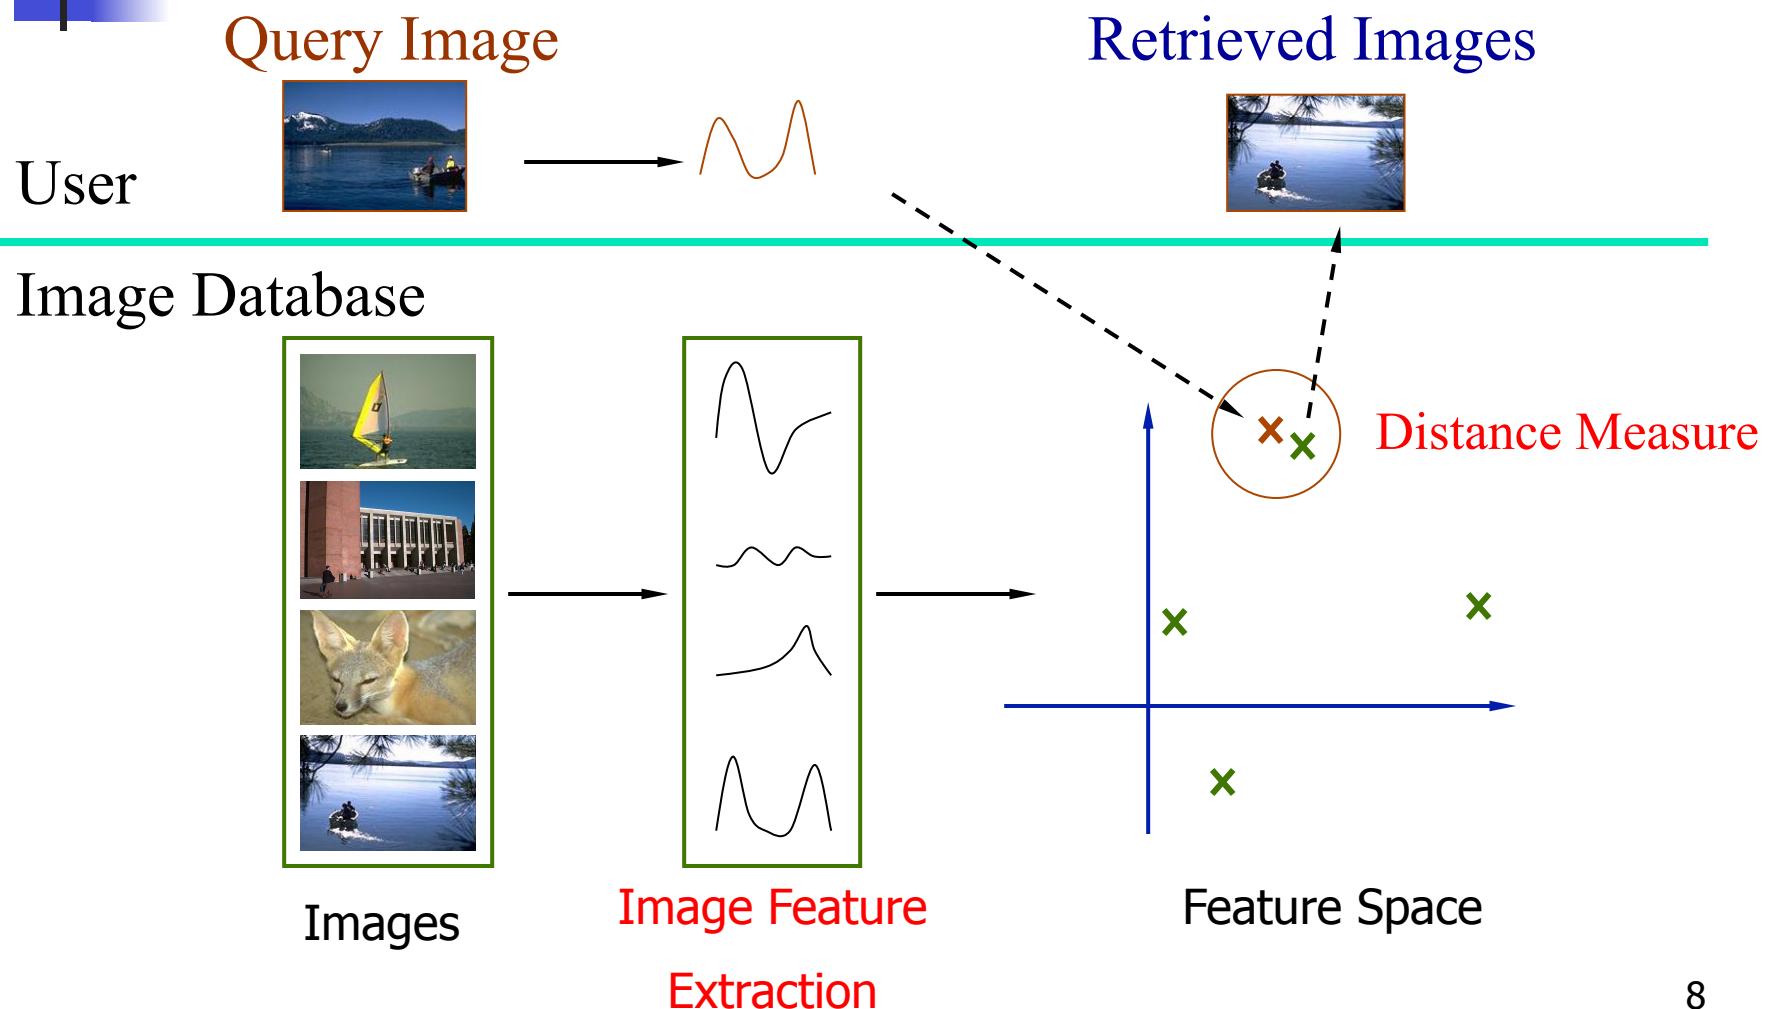
<http://www.informedia.cs.cmu.edu/>

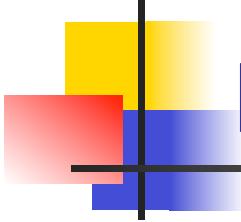
- Started in 1994 and active till 2007
- The largest academia research project on image retrieval

Google Image search:

https://www.google.com/imghp?gws_rd=ssl

Image Features / Distance Measures

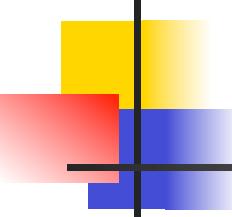




Features

- Color (histograms, gridded layout, wavelets)
- Texture (Laws, Gabor filters, local binary partition)
- Shape (first segment the image, then use statistical or structural shape similarity measures)
- Objects and their relationships

This is the most powerful, but you have to be able to recognize the objects!



Color Histograms

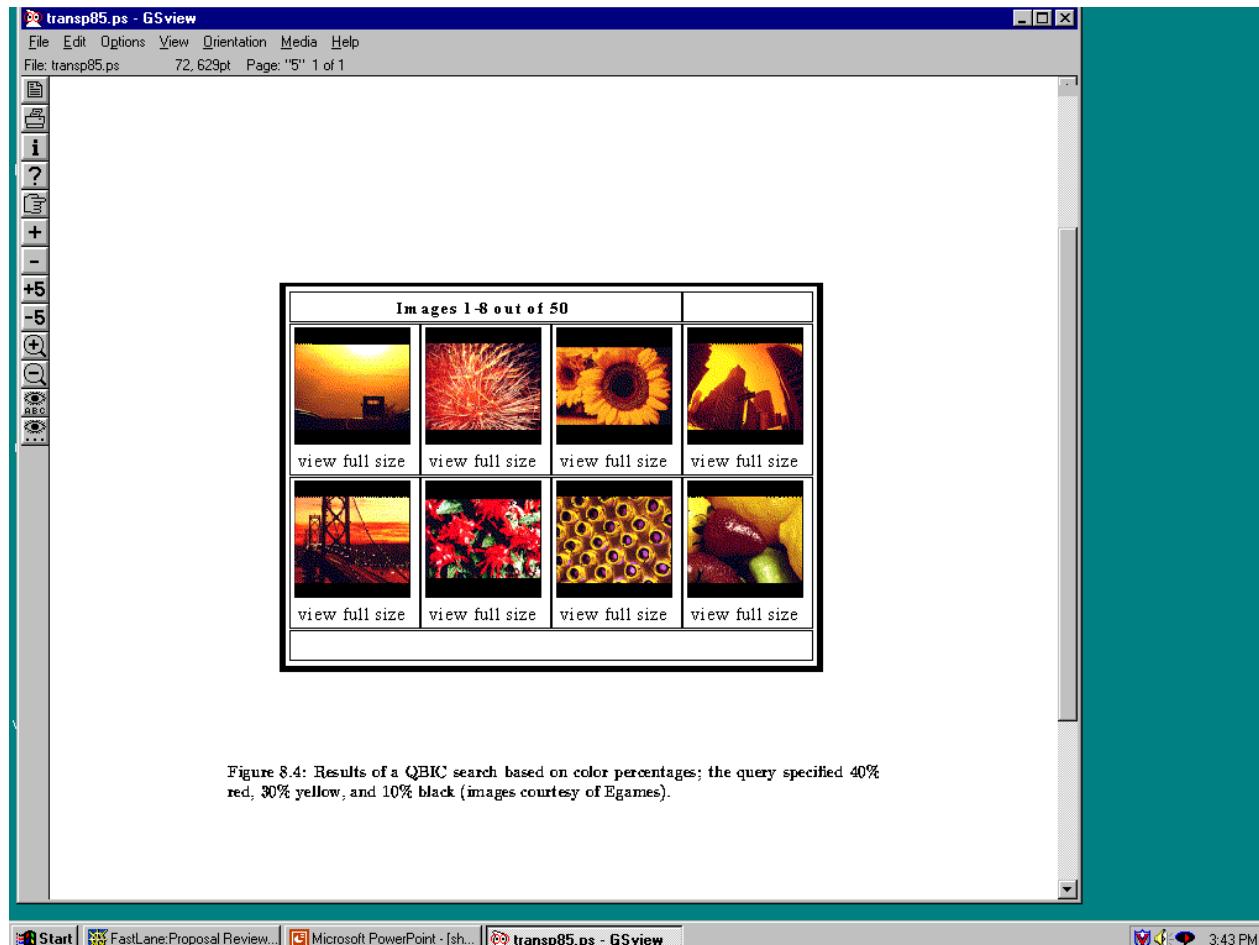
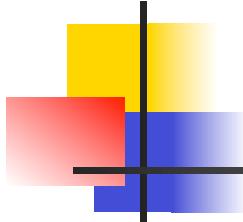


Figure 8.4: Results of a QBIC search based on color percentages; the query specified 40% red, 30% yellow, and 10% black (images courtesy of Egames).

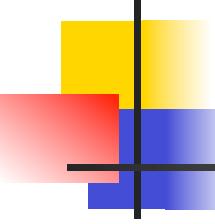


QBIC's Histogram Similarity

The QBIC color histogram distance is:

$$d_{hist}(I, Q) = (h(I) - h(Q))^T \mathbf{A} (h(I) - h(Q))$$

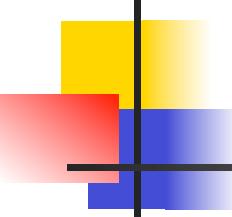
- $h(I)$ is a K-bin histogram of a database image
- $h(Q)$ is a K-bin histogram of the query image
- A is a $K \times K$ similarity matrix



Similarity Matrix

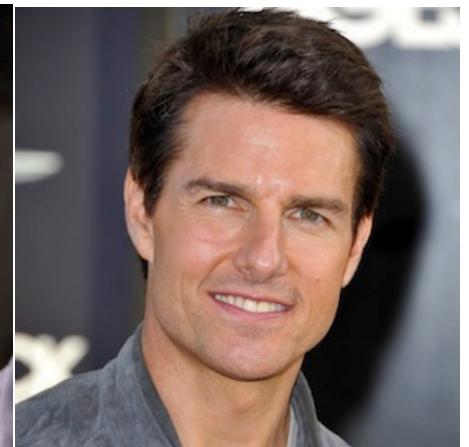
	R	G	B	Y	C	V
R	1	0	0	.5	0	.5
G	0	1	0	.5	.5	0
B	0	0	1		?	
Y				1		
C	?				1	
V						1

How similar is blue to cyan?

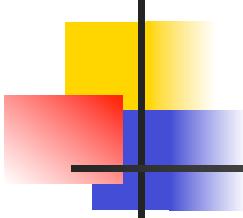


Hair Verification

- Humans use hair for recognition^[1]
- Face is not always suitable for recognition
 - Occlusions due to sunglasses, scarves, etc.
 - Low image resolution/quality
 - Extreme pose variation



Wright and Sladden, "An own gender bias and the importance of hair in face recognition"

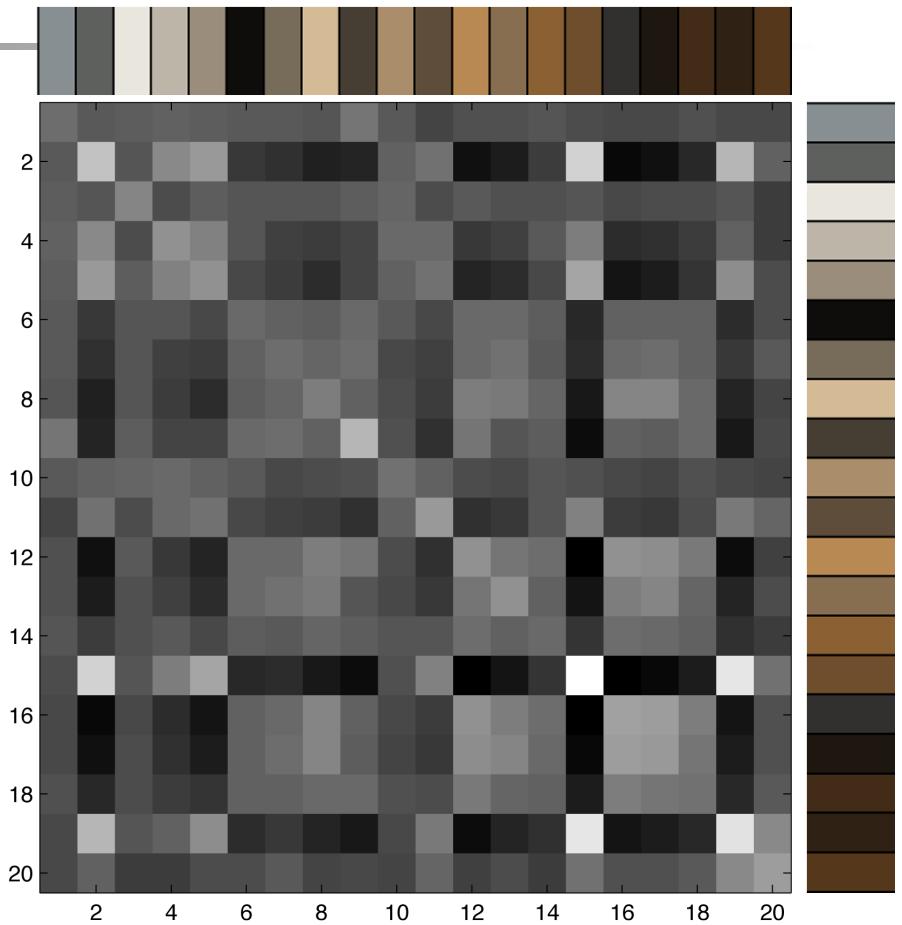


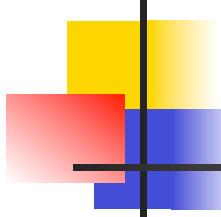
Hair Color

- BoW features
 - Patch sizes: 41, 21
- Cluster in HSV
- Distance Metric Learning

$$\min_A \sum (f_1 - f_2)^T A (f_1 - f_2)$$

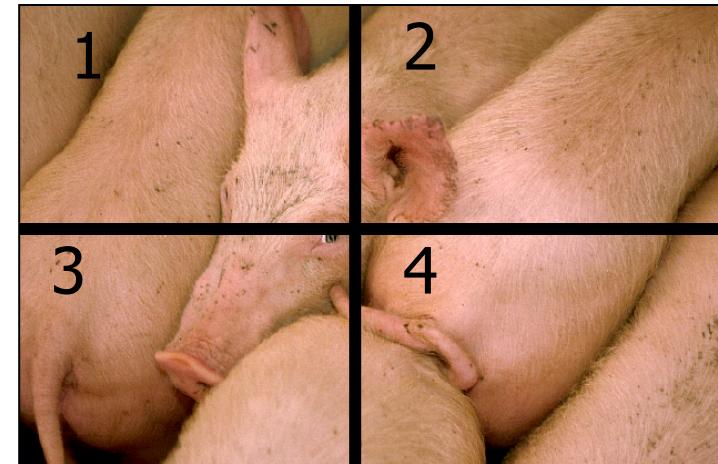
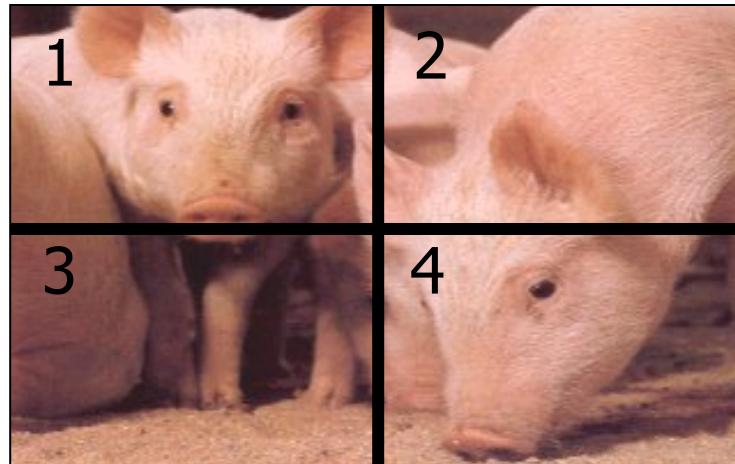
Hist. from image 1 Hist. from image 2





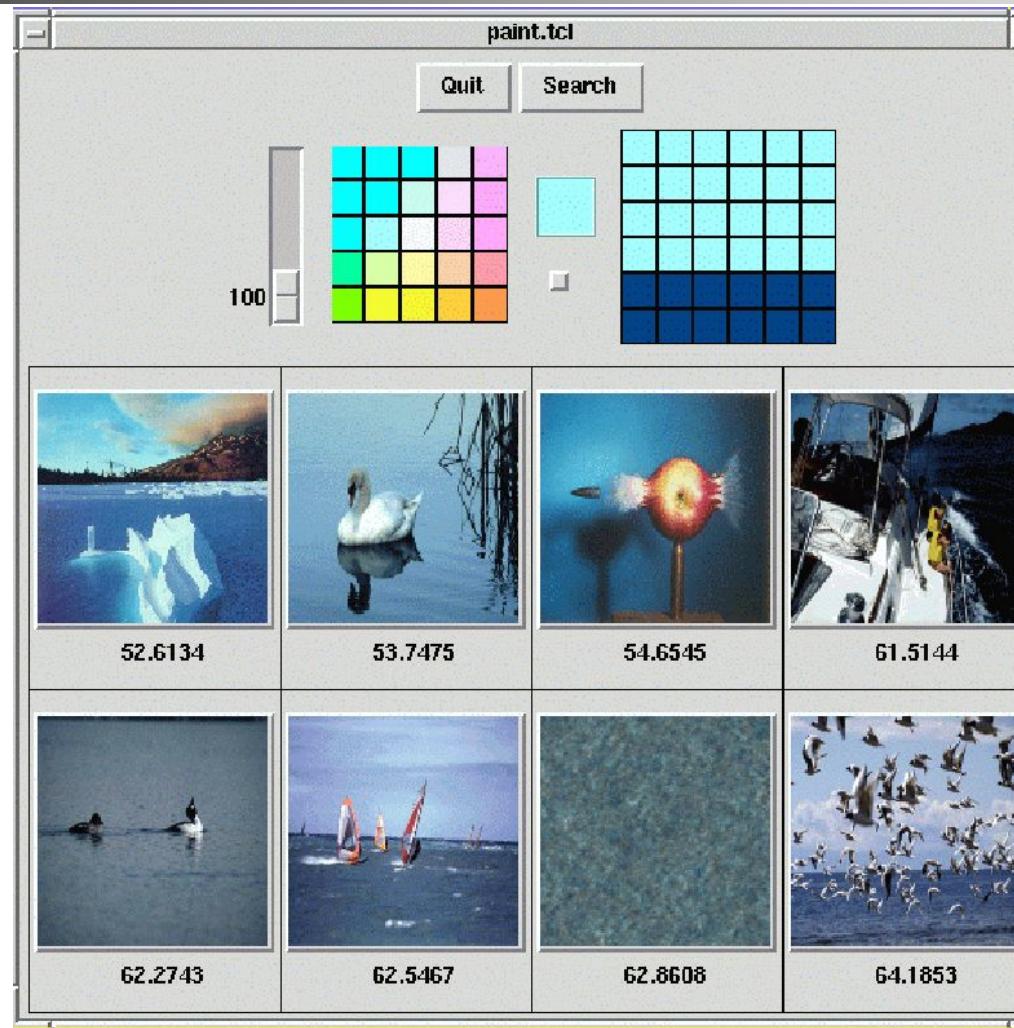
Gridded Color

Gridded color distance is the sum of the color distances in each of the corresponding grid squares.

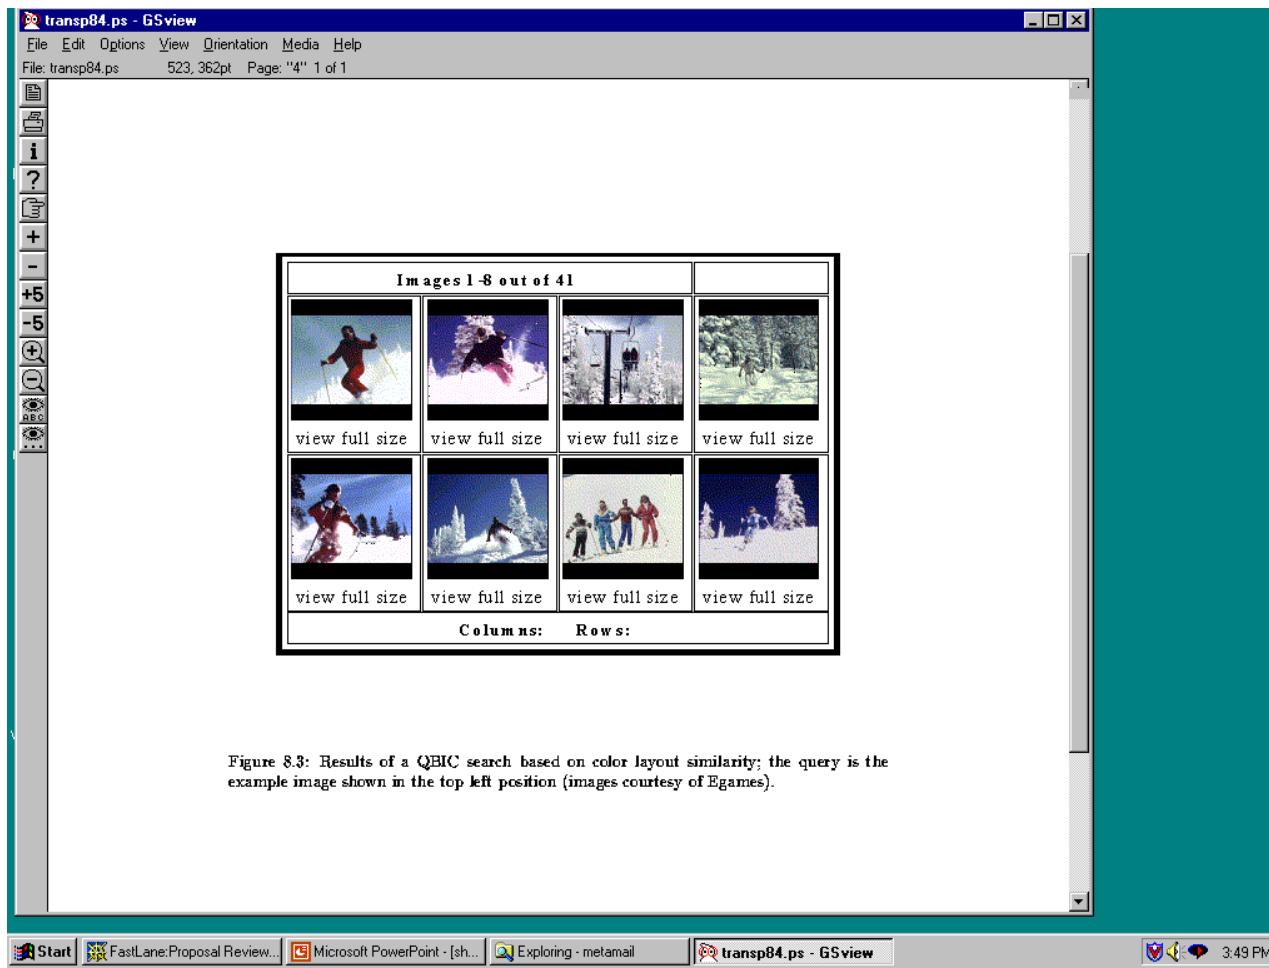


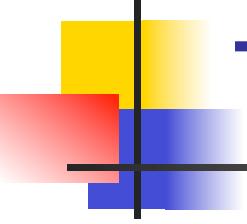
What color distance would you use for a pair of grid squares?

Color Layout



Color Layout (IBM's Gridded Color)





Texture Distances

- Pick and Click (user clicks on a pixel and system retrieves images that have in them a region with similar texture to the region surrounding it).
- Gridded (just like gridded color, but use texture).
- Histogram-based (e.g. compare the LBP histograms).

Laws Texture

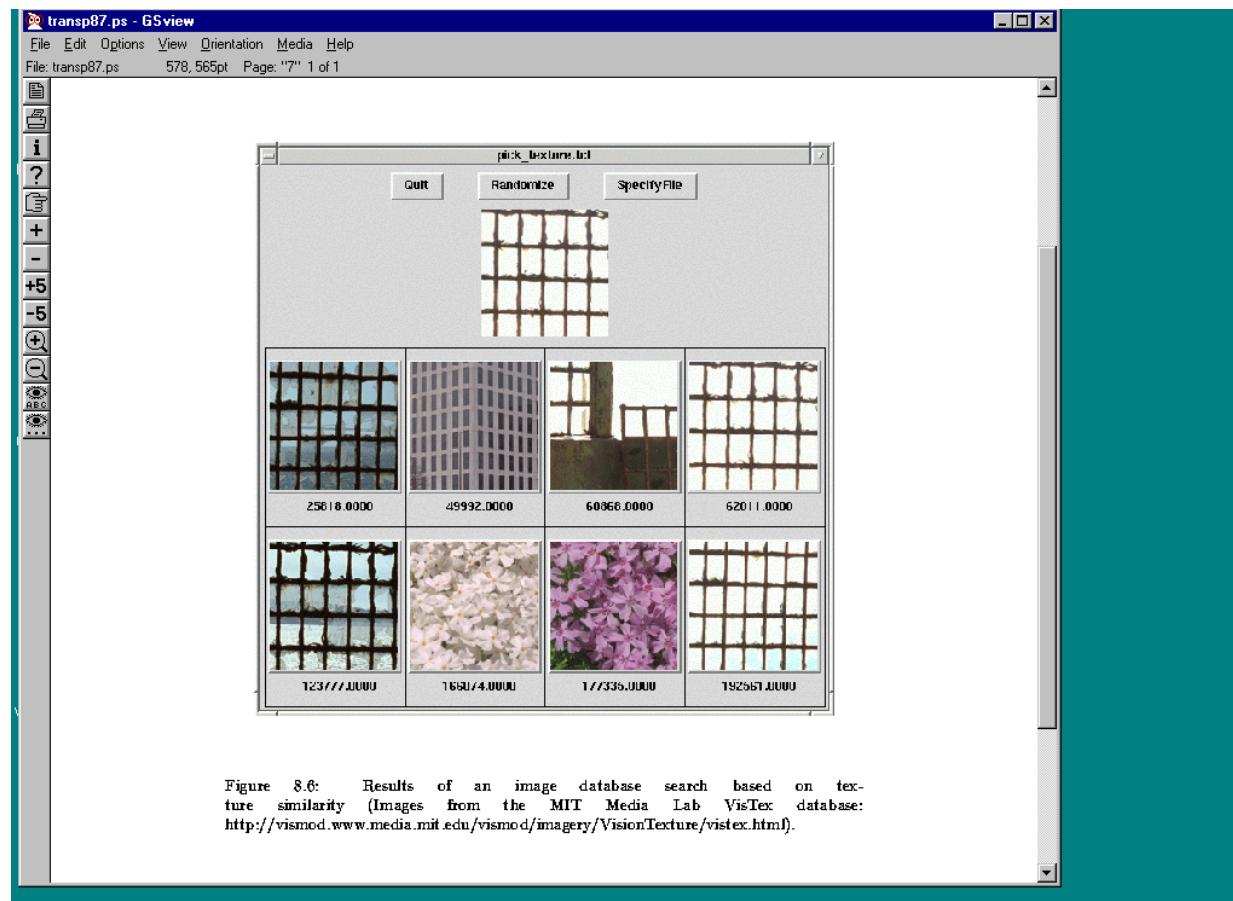
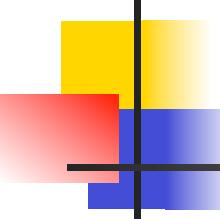


Figure 8.6: Results of an image database search based on texture similarity (Images from the MIT Media Lab VisTex database: <http://vismod.www.media.mit.edu/vismod/imagery/VisionTexture/vistex.html>).

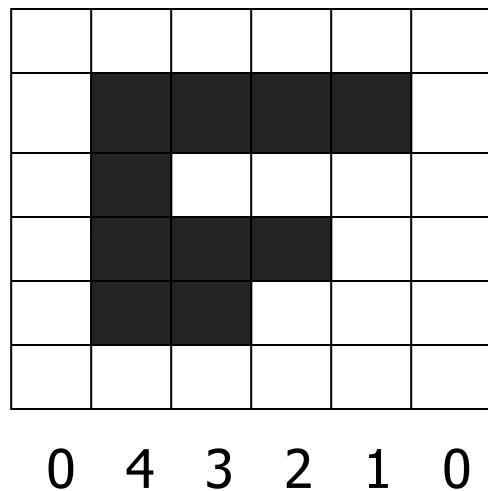
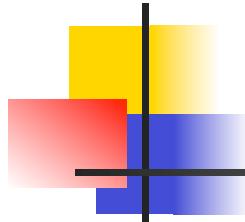
L5 = [+1 +4 6 +4 +1] (Level)
E5 = [-1 -2 0 +2 +1] (Edge)
S5 = [-1 0 2 0 -1] (Spot)
W5 = [-1 +2 0 -2 +1] (Wave)
R5 = [+1 -4 6 -4 +1] (Ripple)



Shape Distances

- Shape goes one step further than color and texture.
- It requires identification of regions to compare.
- There have been many shape similarity measures suggested for pattern recognition that can be used to construct shape distance measures.

Global Shape Properties: Projection Matching



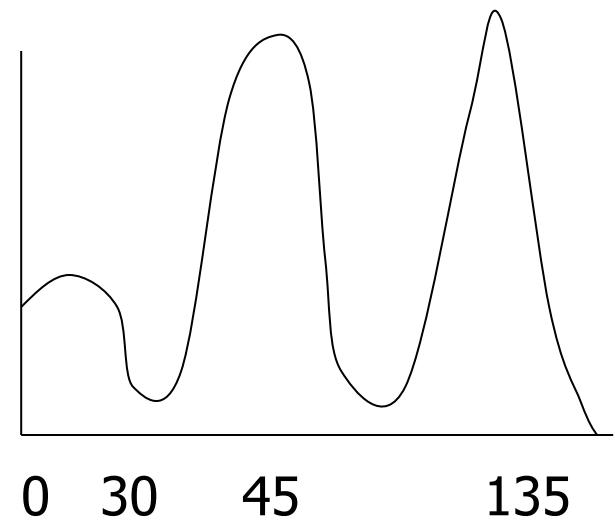
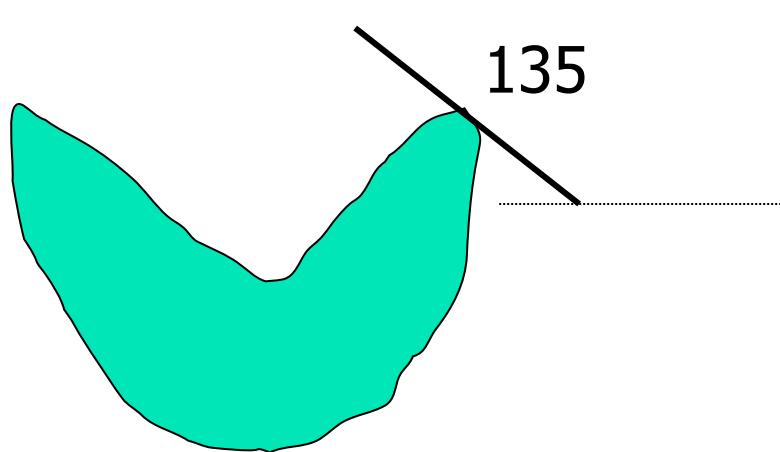
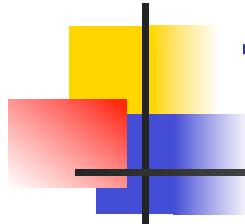
0
4
1
3
2
0

→ Feature Vector
 $(0,4,1,3,2,0,0,4,3,2,1,0)$

In projection matching, the horizontal and vertical projections form a histogram.

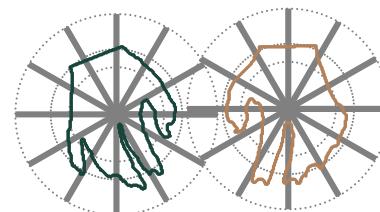
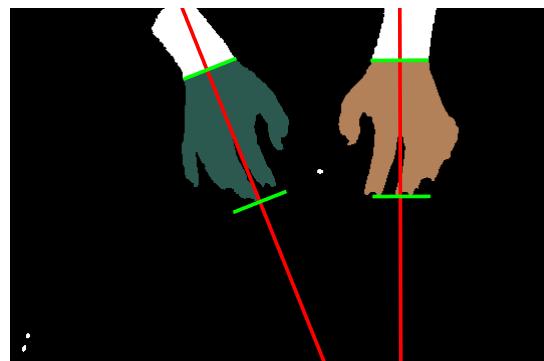
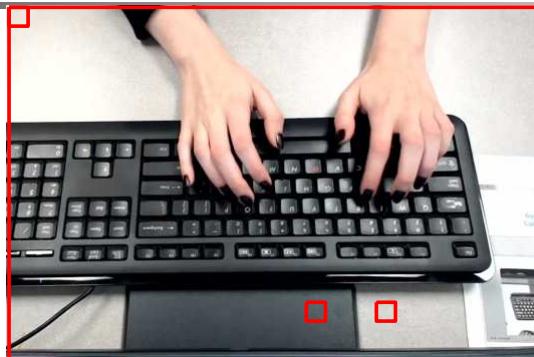
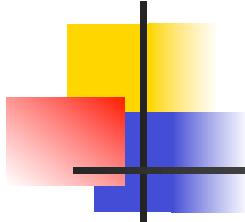
What are the weaknesses of this method? strengths?

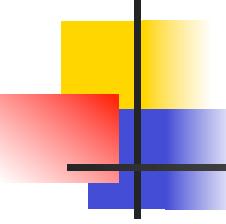
Global Shape Properties: Tangent-Angle Histograms



Is this feature invariant to starting point?

Global Shape Properties: Shape Context



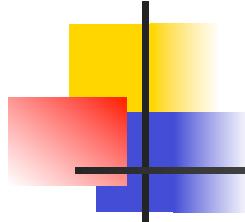


Boundary Matching

- Fourier Descriptors
- Elastic Matching

The distance between query shape and image shape has two components:

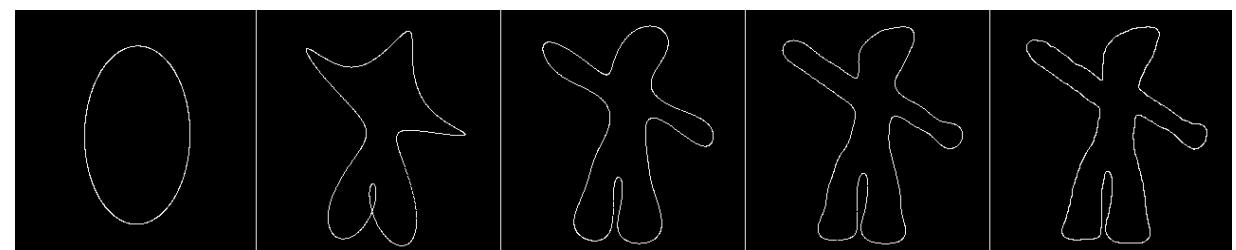
1. energy required to deform the query shape into one that best matches the image shape
2. a measure of how well the deformed query matches the image



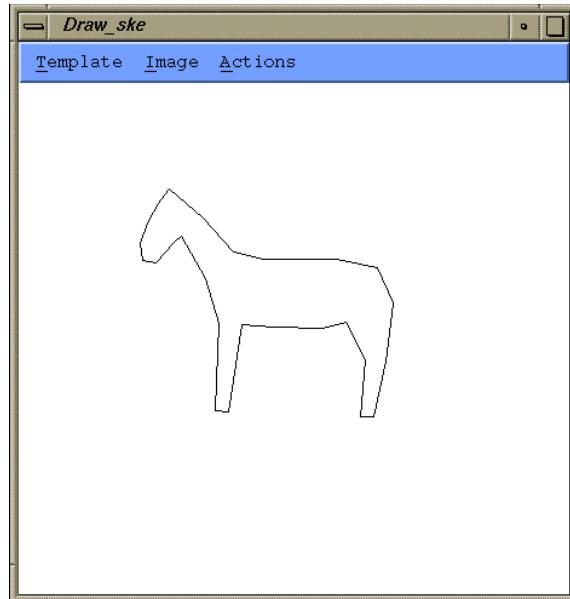
Fourier Descriptors

$$z[m] = x[m] + jy[m]$$

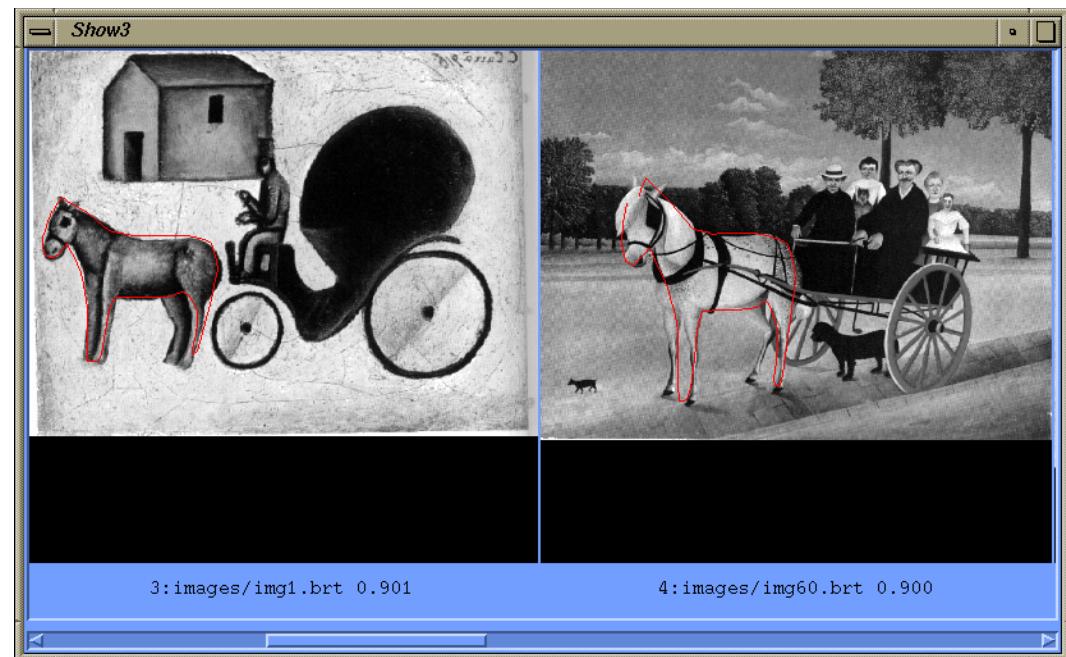
$$Z[k] = DFT[z[m]] = \frac{1}{N} \sum_{m=0}^{N-1} z[m] e^{-j2\pi mk/N} \quad (k = 0, \dots, N-1)$$



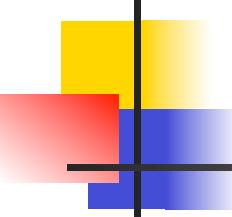
Del Bimbo Elastic Shape Matching



query



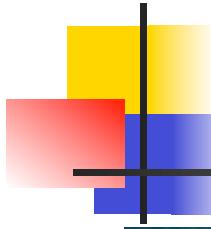
retrieved images



Regions and Relationships

- Segment the image into **regions**
- Find their **properties** and **interrelationships**
- Construct a **graph** representation with nodes for regions and edges for spatial relationships
- Use **graph matching** to compare images

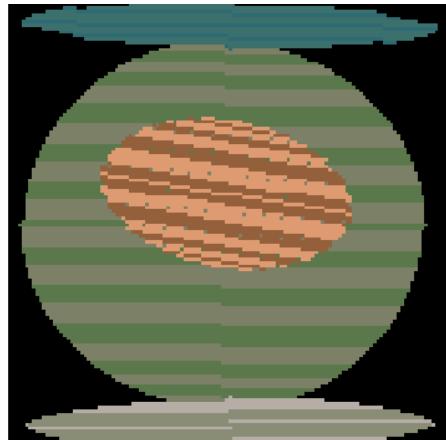
Like
what?



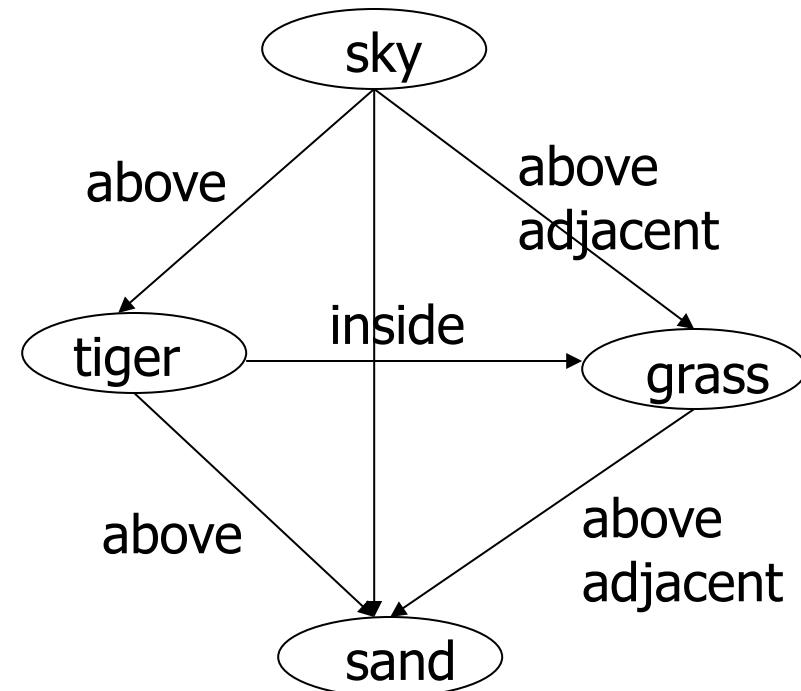
Tiger Image as a Graph



image



abstract regions

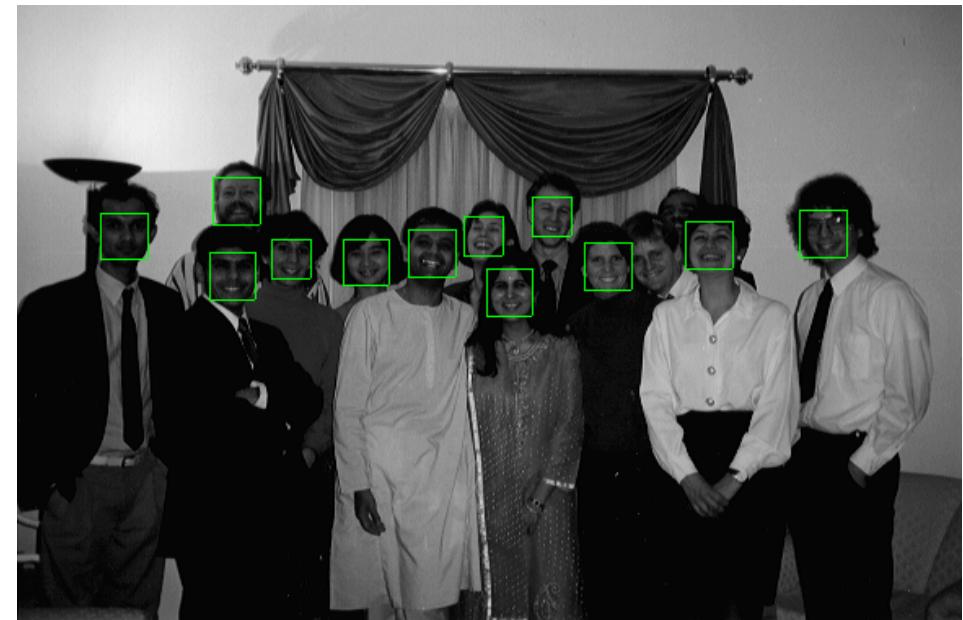


Object Detection: Rowley's Face Finder

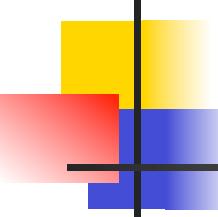
- 1. convert to gray scale
- 2. normalize for lighting*
- 3. histogram equalization
- 4. apply neural net(s)
trained on 16K images

What data is fed to
the classifier?

32 x 32 windows in
a pyramid structure



* Like first step in Laws algorithm, p. 220



Fleck and Forsyth's Flesh Detector

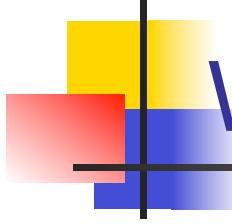
The “Finding Naked People” Paper

- Convert RGB to HSI
- Use the intensity component to compute a texture map
$$\text{texture} = \text{med2}(|I - \text{med1}(I)|)$$
 median filters of
radii 4 and 6
- If a pixel falls into either of the following ranges,
it's a potential skin pixel

$\text{texture} < 5, 110 < \text{hue} < 150, 20 < \text{saturation} < 60$

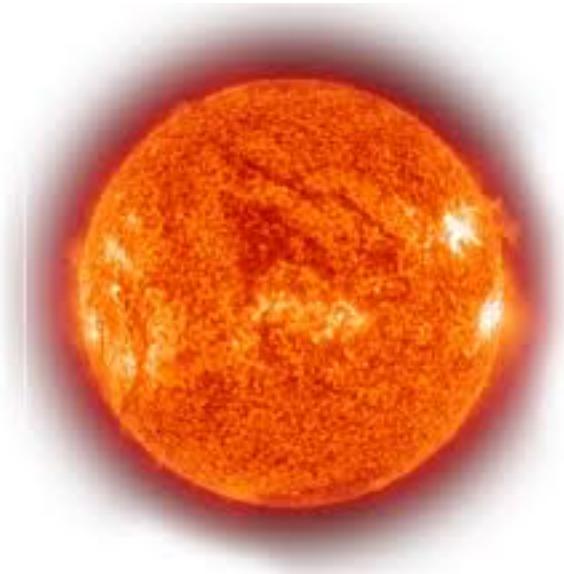
$\text{texture} < 5, 130 < \text{hue} < 170, 30 < \text{saturation} < 130$

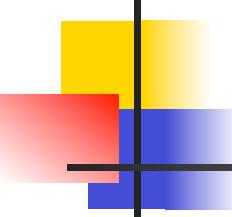
Look for **LARGE** areas that satisfy this to identify pornography.



Weakness of Low-level Features

- Can't capture the high-level concepts



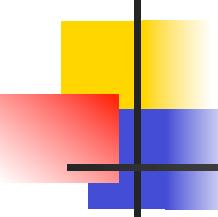


Relevance Feedback

In real interactive CBIR systems, the user should be allowed to interact with the system to “refine” the results of a query until he/she is satisfied.

Relevance feedback work has been done by a number of research groups, e.g.

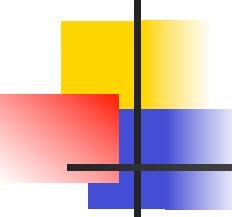
- The Photobook Project (Media Lab, MIT)
- The Leiden Portrait Retrieval Project
- The MARS Project (Tom Huang’s group at Illinois)



Information Retrieval Model*

- An IR model consists of:
 - a document model
 - a query model
 - a model for computing similarity between documents and the queries
- Term (keyword) weighting
- Relevance Feedback

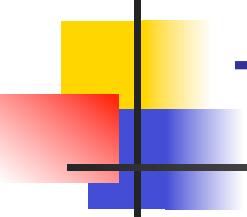
*from Rui, Huang, and Mehrotra's work



Term weighting

- Term weight
 - assigning different weights for different keyword(terms) according their relative importance to the document
- define w_{ik} to be the weight for term t_k , $k=1,2,\dots,N$, in the document i
- document i can be represented as a weight vector in the term space

$$D_i = [w_{i1}; w_{i2}; \dots; w_{iN}]$$



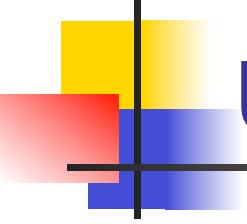
Term weighting

- The query Q also is a weight vector in the term space

$$Q = [w_{q1}; w_{q2}; \dots; w_{qN}]$$

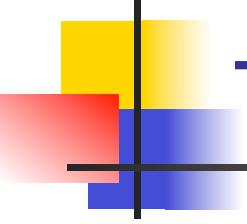
- The similarity between D and Q

$$\text{Sim}(D, Q) = \frac{D \cdot Q}{\|D\| \|Q\|}$$



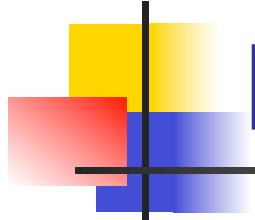
Using Relevance Feedback

- The CBIR system should automatically adjust the weights that were given by the user for the relevance of previously retrieved documents
- Most systems use a statistical method for adjusting the weights.



The Idea of Gaussian Normalization

- If all the relevant images have **similar** values for component j
 - the component j is **relevant** to the query
- If all the relevant images have very **different** values for component j
 - the component j is **not relevant** to the query
- the inverse of the standard deviation of the related image sequence is a good measure of the weight for component j
- **the smaller the variance, the larger the weight**

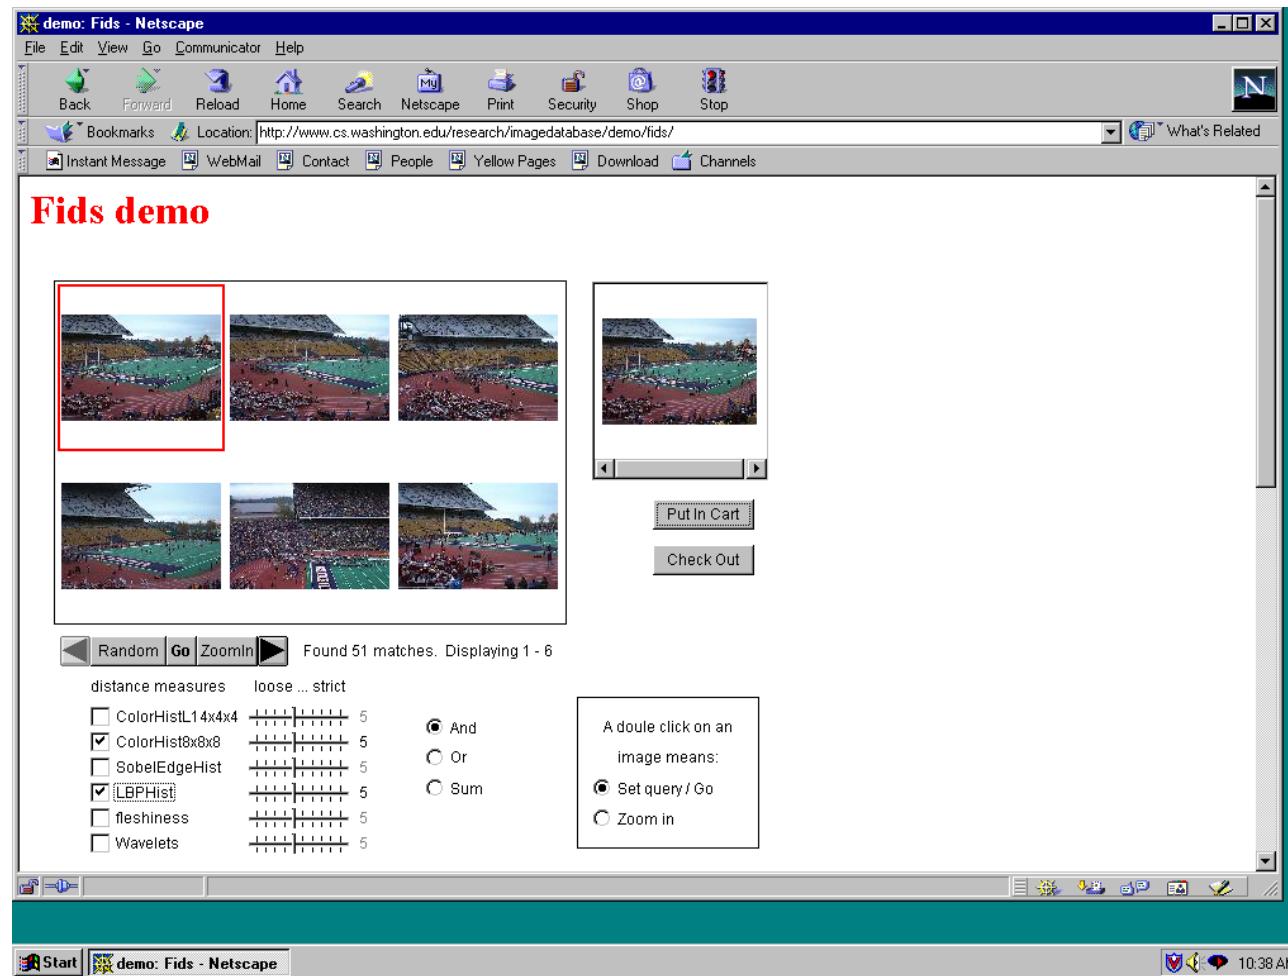


Relevance Feedback Example

<http://www.ifp.illinois.edu/~xzhou2/demo/cbir.html>

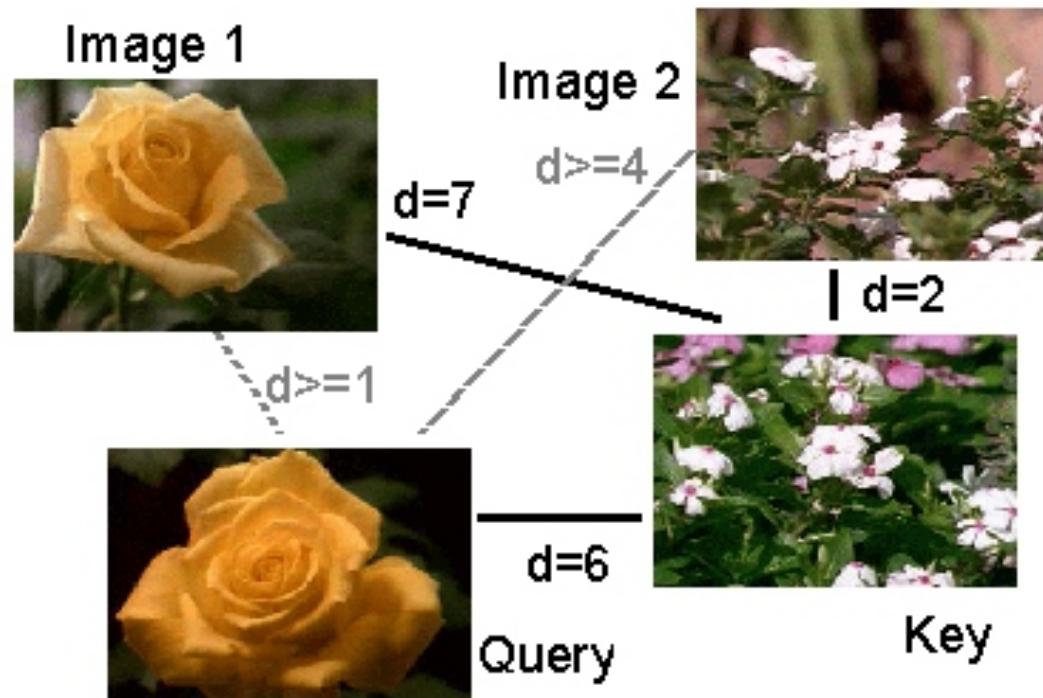
Andy Berman's FIDS System

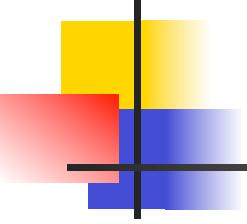
multiple distance measures
Boolean and linear combinations
efficient indexing using images as keys



Andy Berman's FIDS System:

Use of **key images** and the triangle inequality
for efficient retrieval.





Andy Berman's FIDS System:

Bare-Bones Triangle Inequality Algorithm

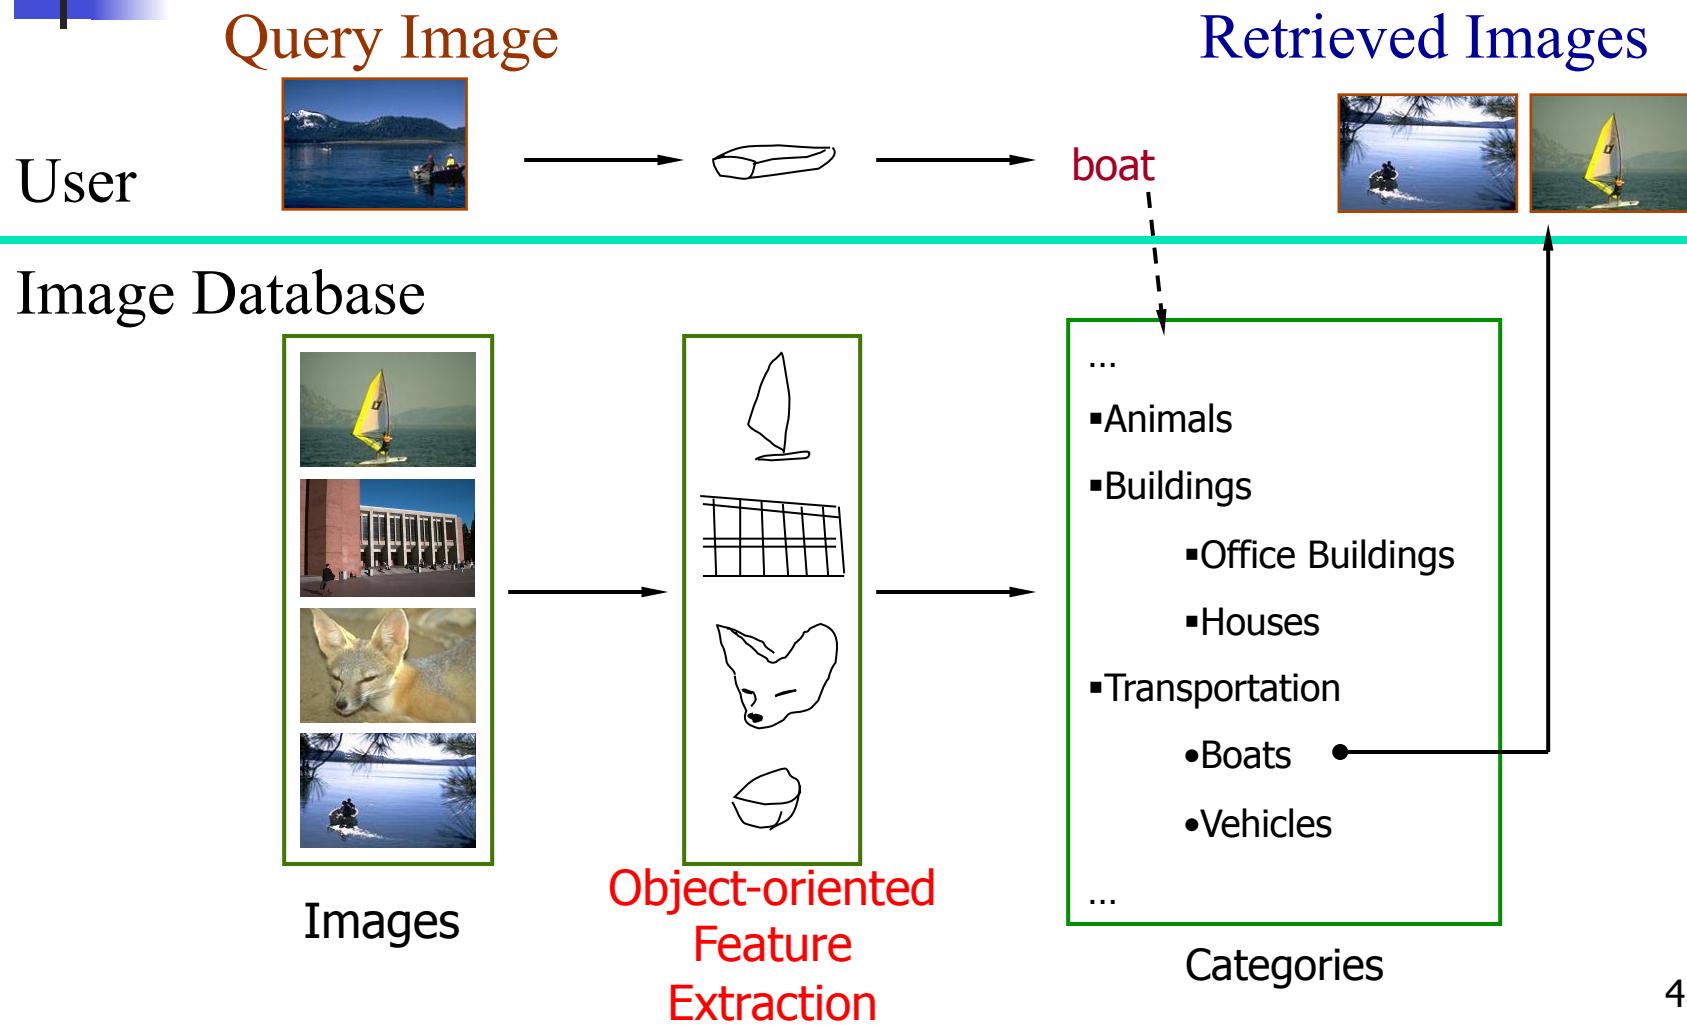
Offline

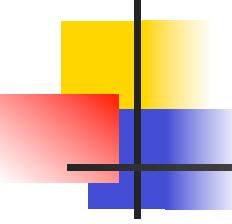
1. Choose a small set of key images
2. Store distances from database images to keys

Online (given query Q)

1. Compute the distance from Q to each key
2. Obtain lower bounds on distances to database images
3. Threshold or return all images in order of lower bounds

Current Research Objective





Overall Approach

- Develop object recognizers for common objects
- Use these recognizers to design a new set of both low- and high-level features
- Design a learning system that can use these features to recognize classes of objects

Boat Recognition

demo: boat recognition - Netscape

File Edit View Go Communicator Help

Bookmarks Location: http://www.cs.washington.edu/research/imagedatabase/demo/boat/ What's Related

Instant Message WebMail Contact People Yellow Pages Download Channels

Boat Recognition

1. Select an image: boat/Q7180237.jpg 2. Select a processor: OR_sailboat 3. Click process>>

Options:

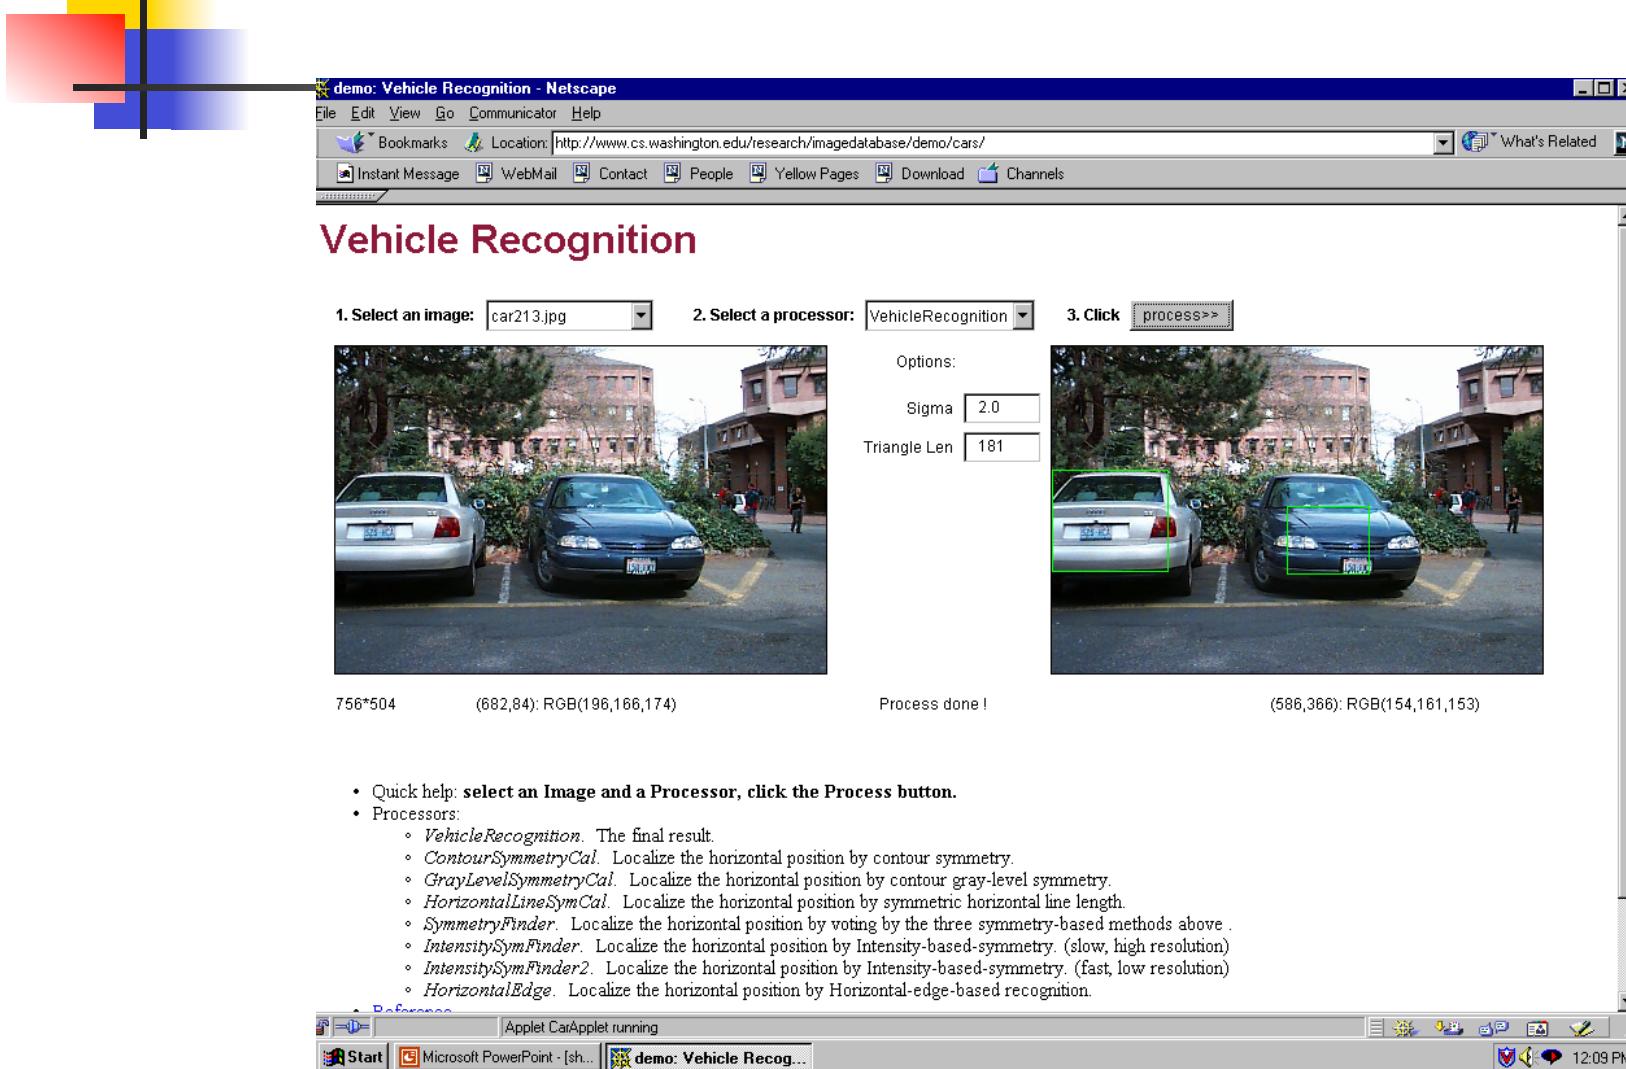
320*240 Process done! (300,12): RGB(0,0,0)

- Quick help: select an Image and a Processor, click the Process button.
- Processors:
 - OR_sky. Sky recognition
 - OR_sea. Sea recognition
 - OR_boat. Boat recognition
 - OR_sailboat. Sailboat recognition

[comments to yi@cs.washington.edu]
Last Modified: Wednesday, December 31, 1969 16:00:00

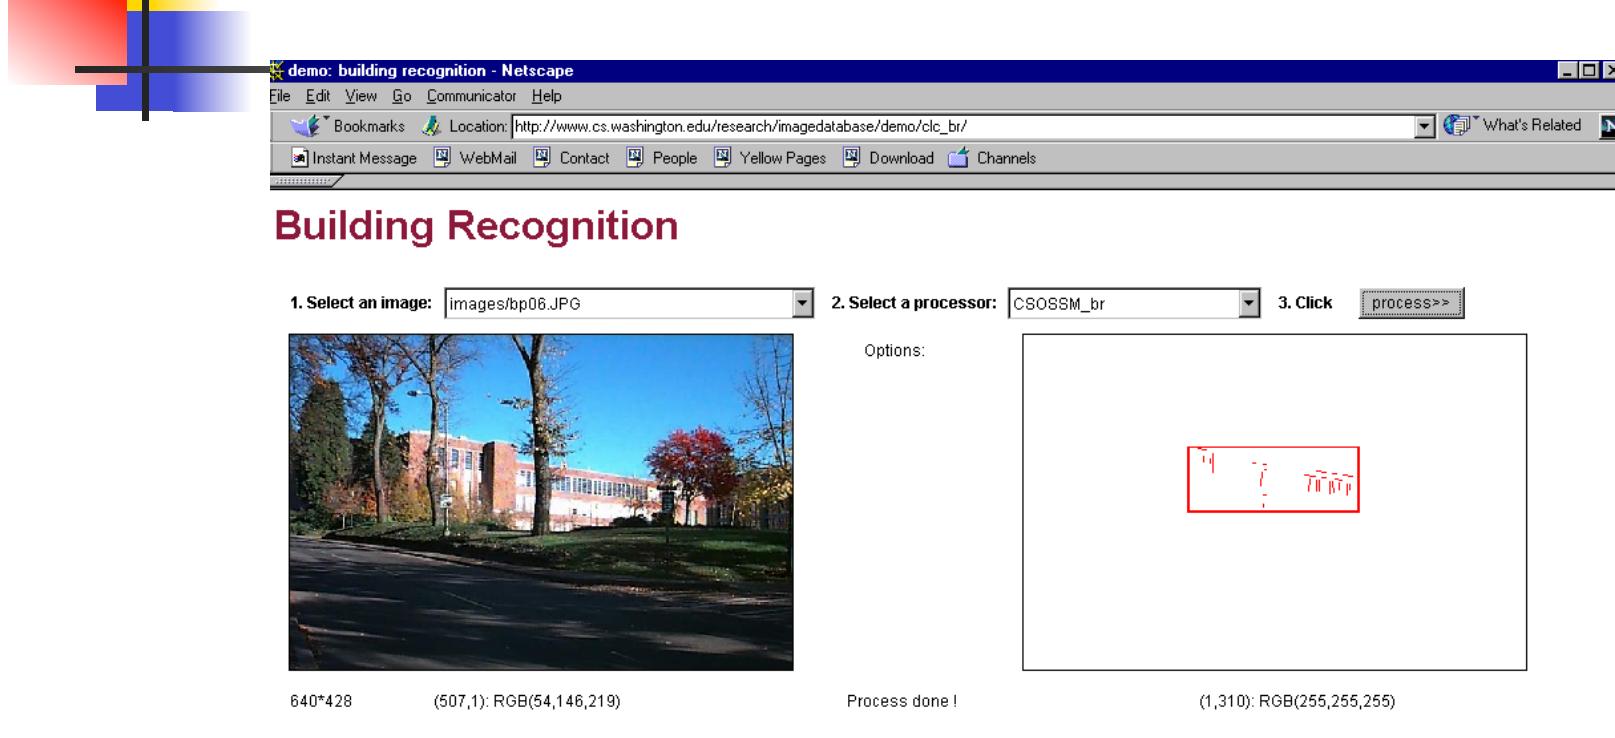
Start Microsoft PowerPoint - sh... demo: boat recognitio... 12:03 PM

Vehicle Recognition



- Quick help: select an Image and a Processor, click the Process button.
- Processors:
 - *VehicleRecognition*. The final result.
 - *ContourSymmetryCal*. Localize the horizontal position by contour symmetry.
 - *GrayLevelSymmetryCal*. Localize the horizontal position by contour gray-level symmetry.
 - *HorizontalLineSymCal*. Localize the horizontal position by symmetric horizontal line length.
 - *SymmetryFinder*. Localize the horizontal position by voting by the three symmetry-based methods above .
 - *IntensitySymFinder*. Localize the horizontal position by Intensity-based-symmetry. (slow, high resolution)
 - *IntensitySymFinder2*. Localize the horizontal position by Intensity-based-symmetry. (fast, low resolution)
 - *HorizontalEdge*. Localize the horizontal position by Horizontal-edge-based recognition.

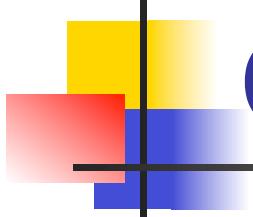
Building Recognition



- Quick help: **select an Image and a Processor, click the Process button.**
- Processors:
 - CSOSSM_br: Building recognition by consistent line clusters

[comments to yi@cs.washington.edu]
Last Modified: Wednesday, December 31, 1969 16:00:00

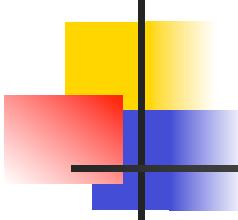




Building Features: Consistent Line Clusters (CLC)

A **Consistent Line Cluster** is a set of lines that are homogeneous in terms of some line features.

- **Color-CLC:** The lines have the same color feature.
- **Orientation-CLC:** The lines are parallel to each other or converge to a common vanishing point.
- **Spatially-CLC:** The lines are in close proximity to each other.

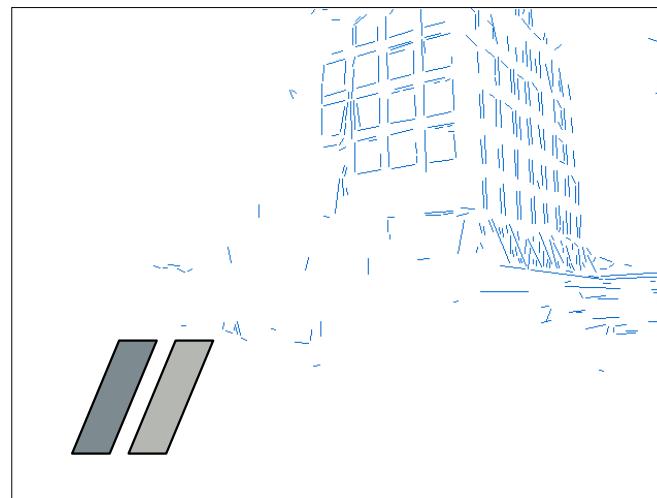


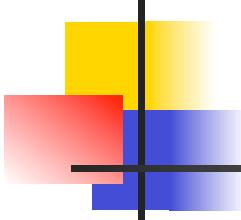
Color-CLC

- Color feature of lines: **color pair** (c_1, c_2)
- Color pair space:
 - RGB $(256^3 * 256^3)$ Too big!
 - Dominant colors $(20 * 20)$
- Finding the color pairs:
 - One line → Several color pairs
- Constructing Color-CLC: **use clustering**



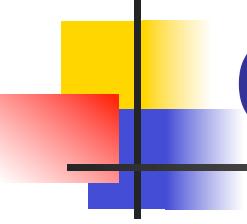
Color-CLC



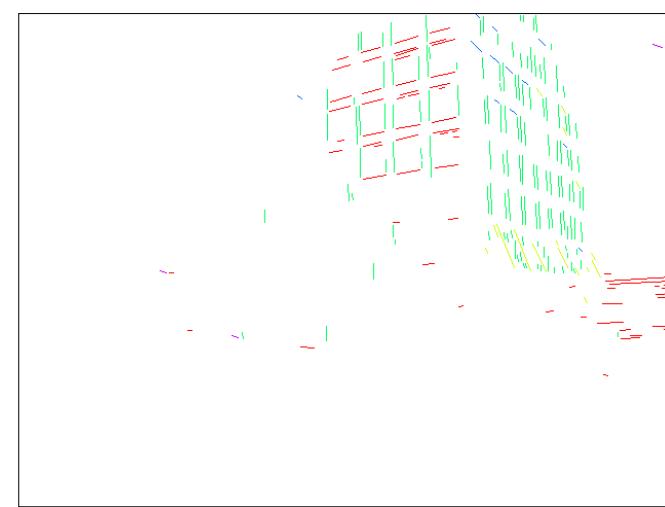
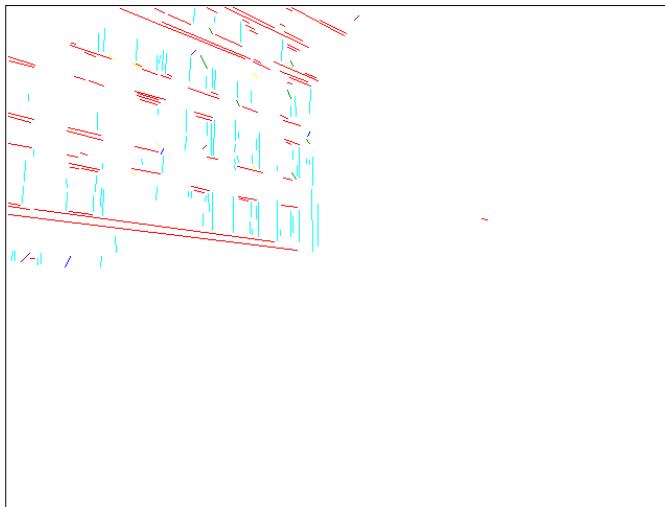


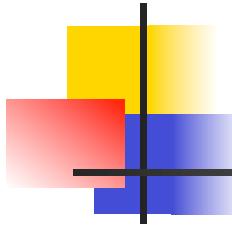
Orientation-CLC

- The lines in an Orientation-CLC are parallel to each other in the 3D world
- The parallel lines of an object in a 2D image can be:
 - Parallel in 2D
 - Converging to a vanishing point (perspective)



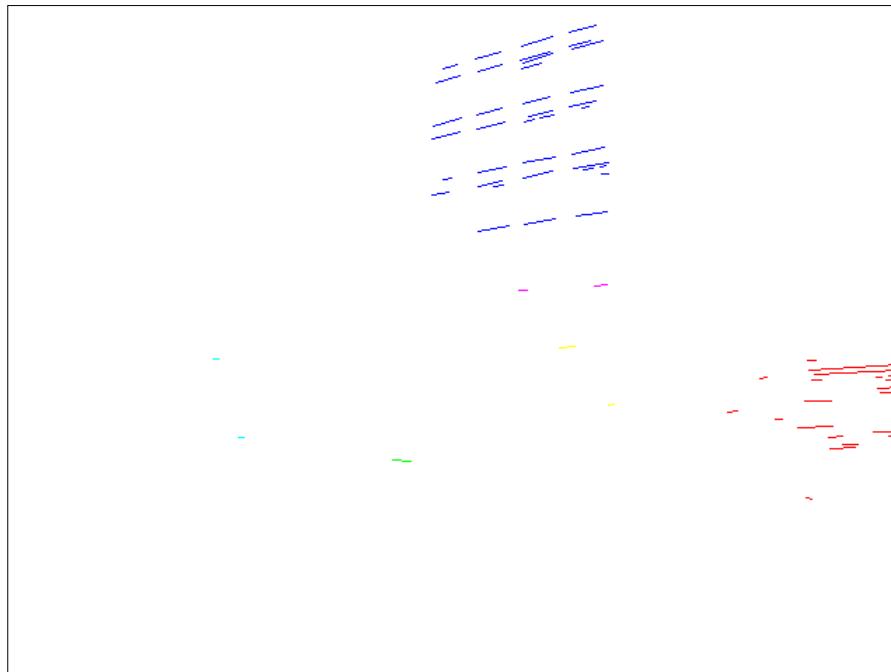
Orientation-CLC

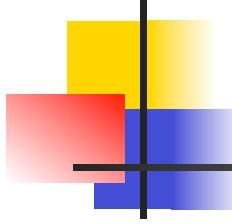




Spatially-CLC

- Vertical position clustering
- Horizontal position clustering

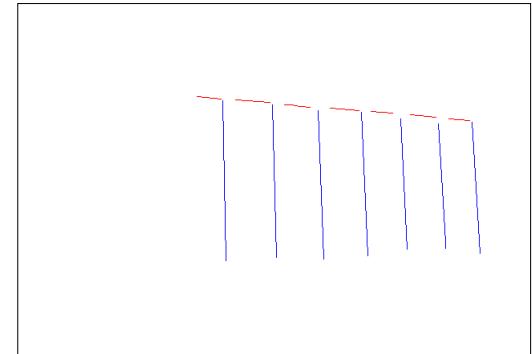
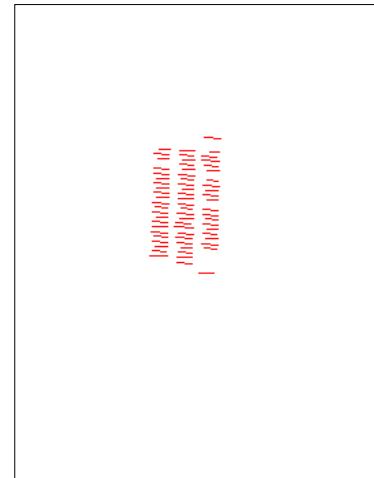


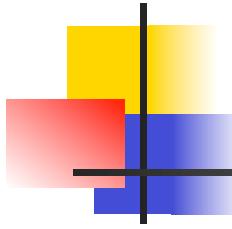


Building Recognition by CLC

Two types of buildings → Two criteria

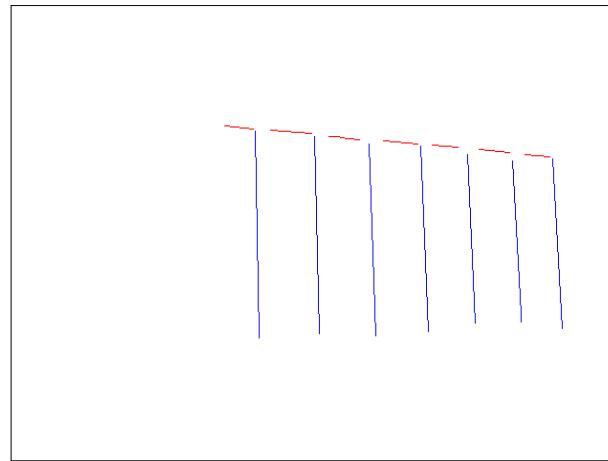
- Inter-relationship criterion
- Intra-relationship criterion





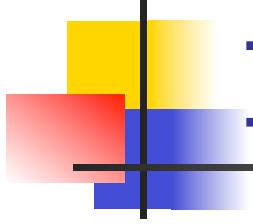
Inter-relationship criterion

$$(N_{c1} > T_{i1} \text{ or } N_{c2} > T_{i1}) \text{ and } (N_{c1} + N_{c2}) > T_{i2}$$



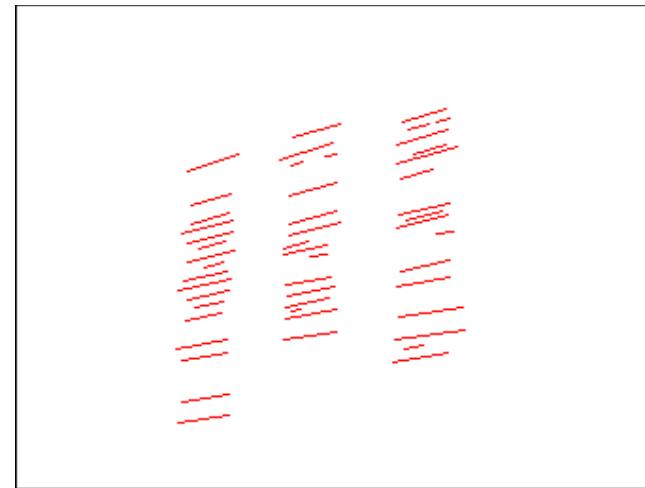
N_{c1} = number of intersecting lines in cluster 1

N_{c2} = number of intersecting lines in cluster 2

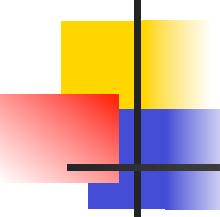


Intra-relationship criterion

$$|S_o| > T_{j1} \text{ or } w(S_o) > T_{j2}$$



S_o = set of heavily overlapping lines in a cluster

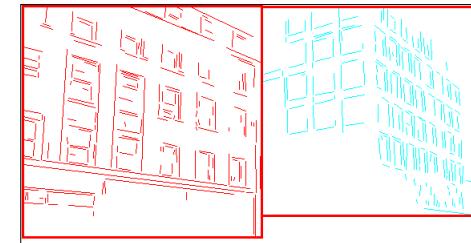
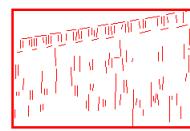
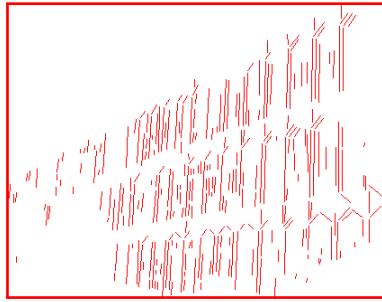


Experimental Evaluation

- Object Recognition
 - 97 well-patterned buildings (bp): **97/97**
 - 44 not well-patterned buildings (bnp): **42/44**
 - 16 not patterned non-buildings (nbnp):
15/16 (one false positive)
 - 25 patterned non-buildings (nbp): **0/25**
- CBIR

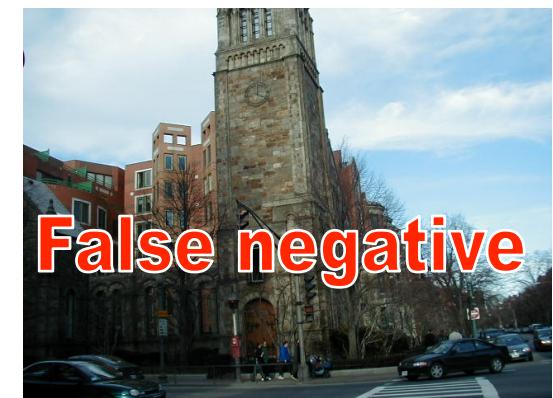
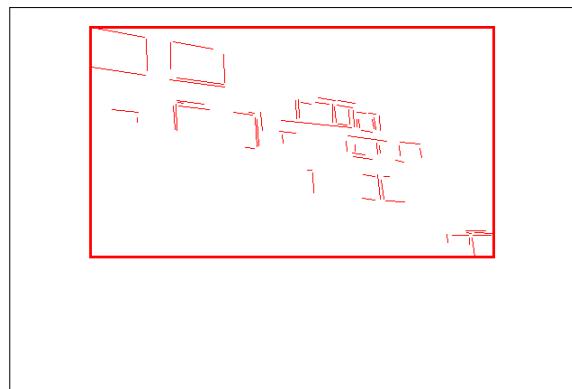
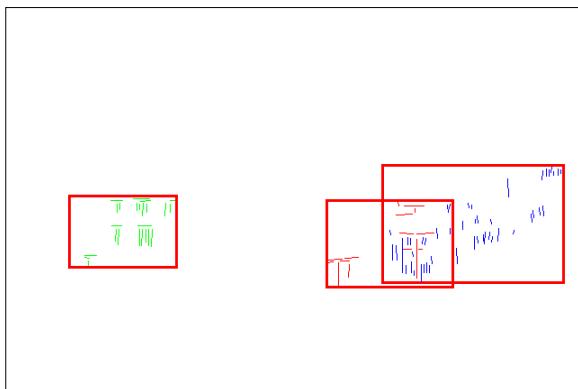
Experimental Evaluation

Well-Patterned Buildings



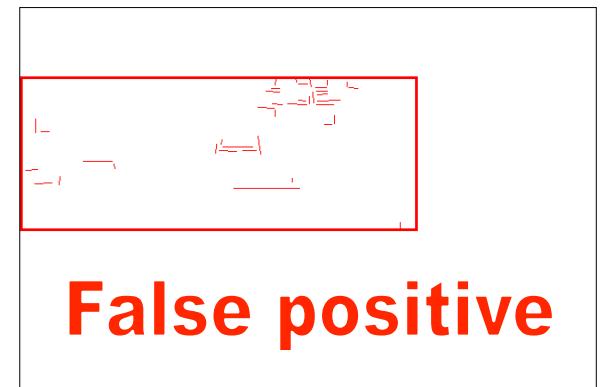
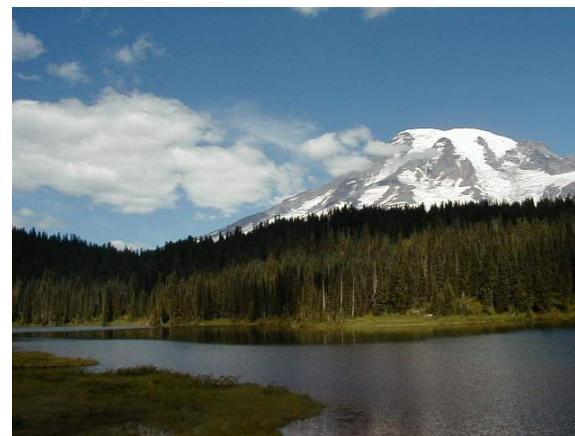
Experimental Evaluation

Non-Well-Patterned Buildings



Experimental Evaluation

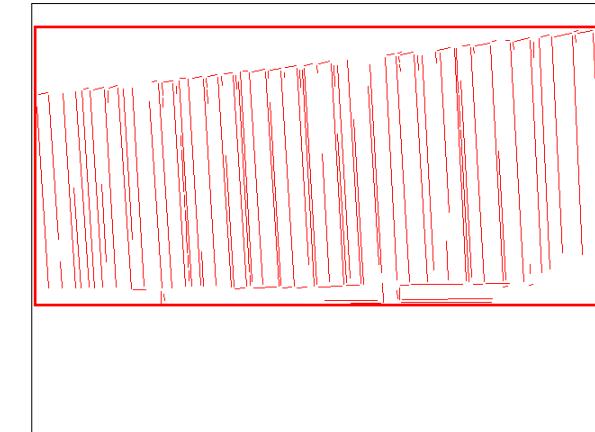
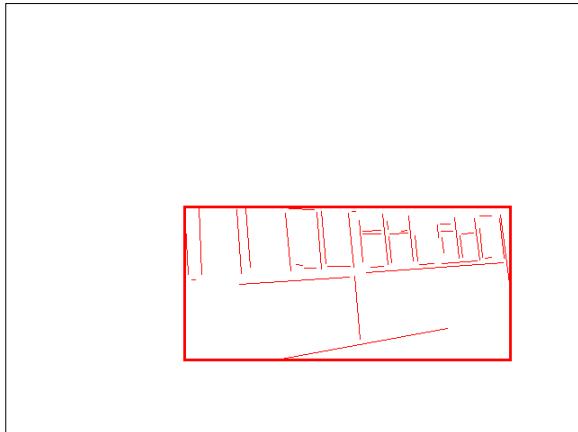
Non-Well-Patterned Non-Buildings



False positive

Experimental Evaluation

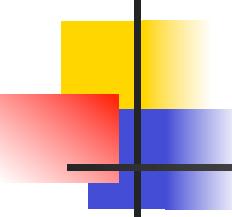
Well-Patterned Non-Buildings (false positives)



Experimental Evaluation (CBIR)

False positives from Yellowstone





Future Work

■ Future Work

- Constructing hierarchically structured clusters
- Using CLC on other objects
- Combining CLC with other features
- Developing a learning approach using hierarchical, multiple classifiers (Chou 2000)

Text Search

- Documents (Web pages) represented by words
- Inverted index links keywords to the documents that contain them
- Keyword query retrieves documents containing that word

Inverted Index Example

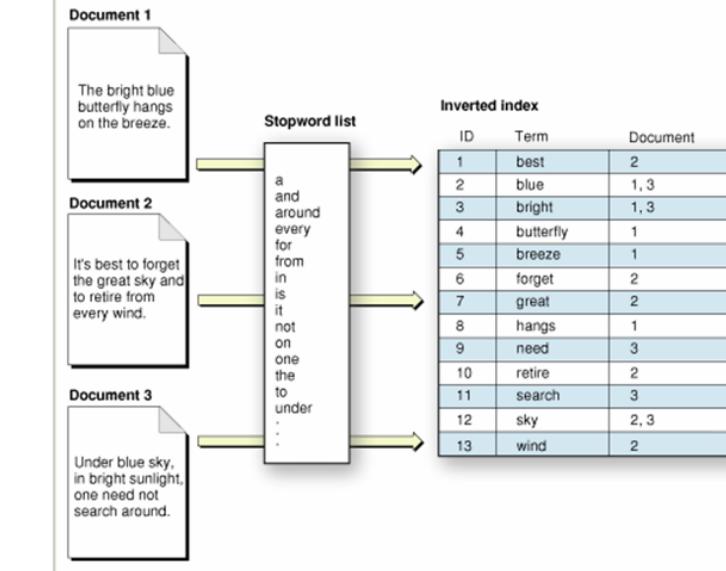
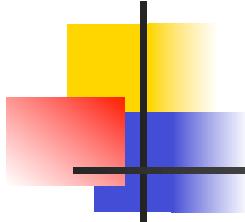


Image from http://developer.apple.com/documentation/UserExperience/Conceptual/SearchKitConcepts/searchKit_basics/chapter_2_section_2.htm

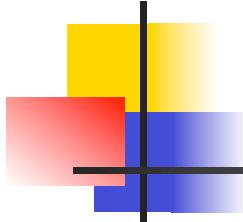
Image Search - Tags



- Search over tags associated with images
 - Users manually add Tags to images
 - FlickR
 - FaceBook
 - Find images with tags that match the query keyword
- Limitations
 - Tags require human effort to create
 - Tags may be wrong



Image Search - Text



- Use text associated with images for search
 - Search web for images
 - Use surrounding text
 - Text in URL for image filename
 - Text in HTML on page
 - Same as text search
- Example: Google Image Search for “Sunset” gives
 - Sunset at Rocky Point in Australia
 - Sunset Beach, Oahu
 - Frank Smiles at Sunset
- Because the keyword “Sunset” was in the title of all these images



Sunset at Rocky Point



Frank Smiles
at Sunset



Sunset Beach

Image Search – Image Descriptors

- SIFT Features
 - (Scale Invariant Feature Transform) Features
 - (2004: David Lowe, UBC)
- Select keypoints regions in image from extrema in scale space
 - Different images have different numbers of keypoints
- Compute feature vectors X for each keypoint region
 - Feature vectors from histogram of gradient directions near the keypoint
 - SIFT features X are 128-dimensional vectors
- Image described by N SIFT features
 - Features are X_1, \dots, X_N
 - N is different for different images



Image Search – Visual Words

- Quantize SIFT features to create “visual words” to represent images
 - (2006: Lienhart, University of Augsburg & Slaney, Yahoo!)
- Cluster SIFT features of representative images
 - Features X are in 128-dimensional space
 - Generate W clusters
 - Clusters define “**visual words**”
 - All features in same cluster are the same “visual word”
- To compute visual words describing an image
 - Compute N SIFT X_1, \dots, X_N features for the image
 - Find nearest cluster center (codeword) to each feature X_j
 - These clusters define the visual words for the image

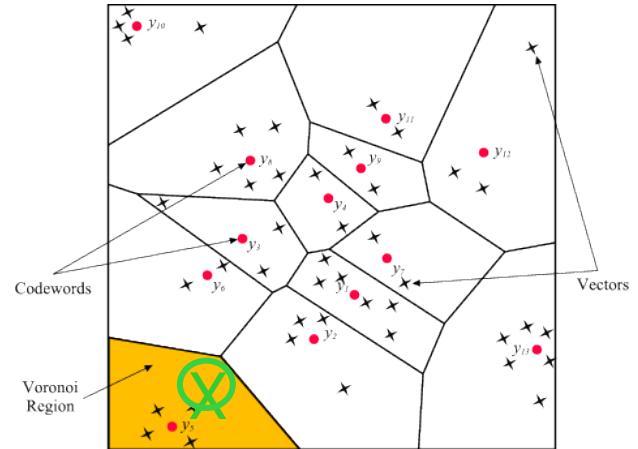
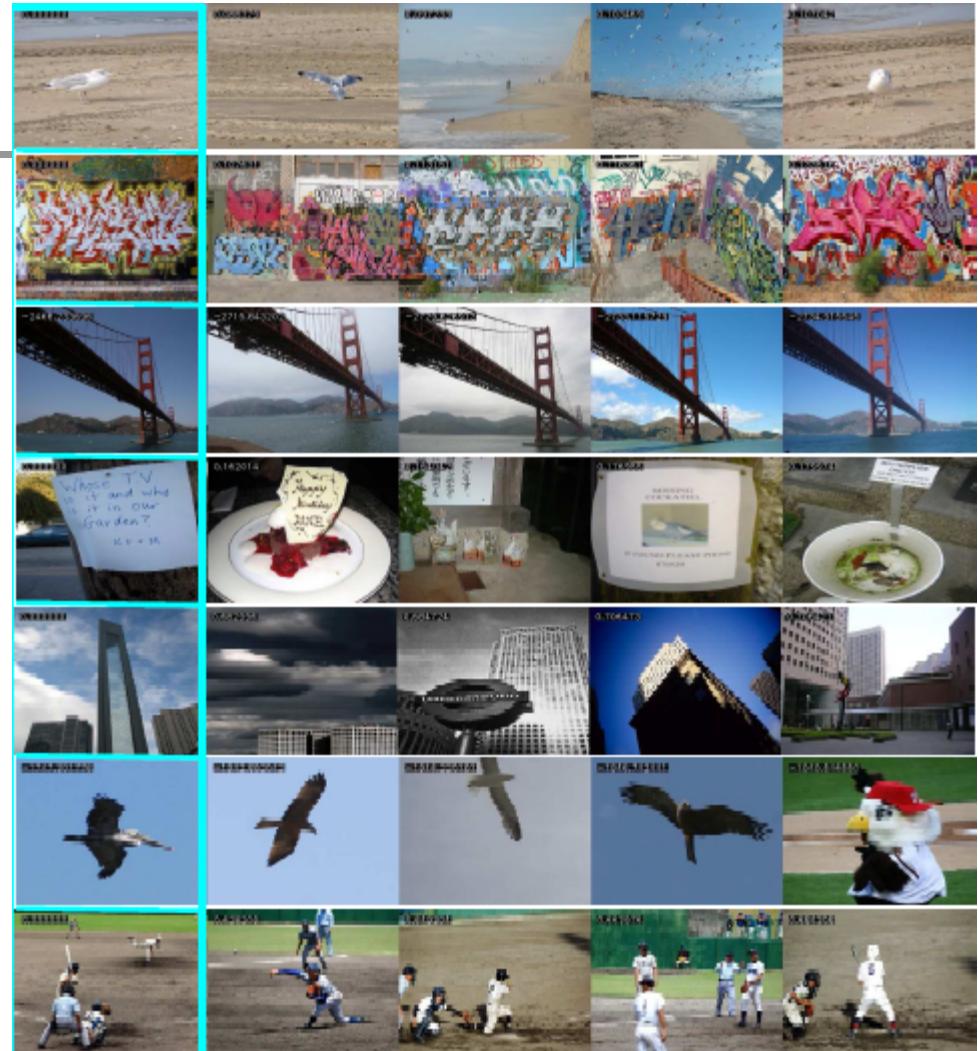


Image Feature 
Visual Word 1 

Image Retrieval – Visual Words

- Image is described by its visual words
 - Just like a document is described by the text words
- Create image index
 - Compute visual words for all images
 - Create a visual word index into the images
- Compute visual words for query image
 - Use query words for retrieval
- Just like text!
 - Except the visual words aren't quite as meaningful

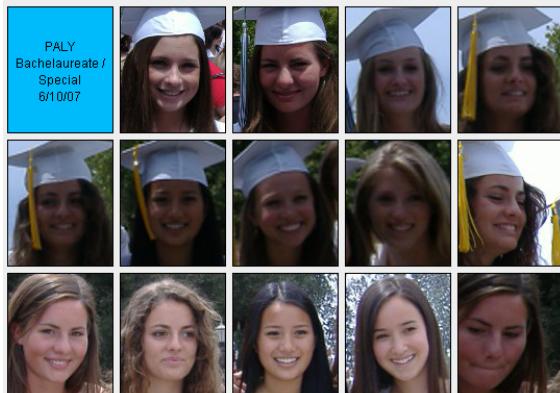


Slide from Lynn Wilcox

Image Search –Faces

Face Detection

- Find faces in images
 - Search for all images with faces
 - Ex: Google advance search for images with faces
 - Good results!
- Example:
- FXPAL Photo Application (2004: Grgenohn et al.)



Faces in Photo Collection



Photo Collection



Slide from Lynn Wilcox Face Detection

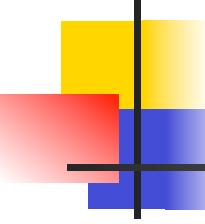


Image Search –Faces

- Face Recognition
 - Search for all images of a particular person
 - Bad results!
- Face Similarity
 - Similarity search based on face features
 - Use face similarity to help manually label faces
 - Good results!
- User Interface for Labeling Faces
 - Drag face to label

